

Zc* v0.61.0

This document is in the Stable state. Assume anything could still change, but limited change should be expected. For more information see: <https://riscv.org/spec-state>

Zc* is a group of extensions which define subsets of the existing C extension (Zca, Zcf) and new extensions which only contain 16-bit encodings.

Zcy* all reuse encodings for c.fld*, c.fsd*.

Table 1. Zc* extension overview

Instruction	Zca	Zcf	Zcb	Zcyh	Zcyp	Zcype	Zcyt
Define a subset of C with the FP load/stores removed							
C excl. c.f*	✓						
The single precision floating point load/stores become a separate extension							
c.flw		✓					
c.flwsp		✓					
c.fsw		✓					
c.fswsp		✓					
Simple operations							
c.zext.b			✓				
c.sext.b			✓				
c.zext.h			✓				
c.sext.h			✓				
c.zext.w			✓				
c.mul			✓				
c.not			✓				
Load/store byte/half							
c.lb				✓			
c.lbu				✓			
c.lh				✓			
c.lhu				✓			
c.sb				✓			
c.sh				✓			
Push/pop and double move							
c.push					✓	✓	
c.pop					✓	✓	
c.popret					✓	✓	

Instruction	Zca	Zcf	Zcb	Zcyh	Zcyp	Zcype	Zcyt
c.popretz					✓	✓	
c.mva01s					✓		
c.mvsa01					✓		
Reserved for EABI versions of push/pop and double move							
c.push.e						✓	
c.pop.e						✓	
c.popret.e						✓	
c.popretz.e						✓	
c.mva01s.e						✓	
c.mvsa01.e						✓	
Table jump							
c.jt							✓
c.jalt							✓

Zca

Zca is all of the existing C extension, *excluding* all 16-bit floating point loads and stores: *c.flw*, *c.flwsp*, *c.fsw*, *c.fswsp*, *c.fld*, *c.fldsp*, *c.fsd*, *c.fsdsp*.

Zcf

Zcf is the existing set of single precision floating point loads and stores: *c.flw*, *c.flwsp*, *c.fsw*, *c.fswsp*.

Zcb

All proposed encodings are currently reserved for all architectures, and have no conflicts with any existing extensions.

The *c.mul* encoding uses the CR register format along with other instructions such as *c.sub*, *c.xor* etc.

NOTE | *c.sext.w* is a pseudo-instruction for *c.addiw rd, 0* (RV64)

RV32	RV64	Mnemonic	Instruction
✓	✓	<i>c.zext.b rsd'</i>	Zero extend byte, 16-bit encoding
✓	✓	<i>c.sext.b rsd'</i>	Sign extend byte, 16-bit encoding
✓	✓	<i>c.zext.h rsd'</i>	Zero extend halfword, 16-bit encoding
✓	✓	<i>c.sext.h rsd'</i>	Sign extend halfword, 16-bit encoding
	✓	<i>c.zext.w rsd'</i>	Zero extend word, 16-bit encoding
✓	✓	<i>c.not rsd'</i>	Bitwise not, 16-bit encoding
✓	✓	<i>c.mul rsd', rs2'</i>	Multiply, 16-bit encoding

Zcyh

This extension reuses some encodings from c.fld, c.fldsp, c.fsd, c.fsdsp. Therefore it is *incompatible* with D. It is fully compatible with F and, with Zdinx and also with [Zca](#).

The instructions are all 16-bit versions of existing 32-bit load/store instructions.

RV32	RV64	Mnemonic	Instruction
✓	✓	c.lbu rd' , uimm($rs1'$)	Load unsigned byte, 16-bit encoding
✓	✓	c.lhu rd' , uimm($rs1'$)	Load unsigned halfword, 16-bit encoding
✓	✓	c.lb rd' , uimm($rs1'$)	Load signed byte, 16-bit encoding
✓	✓	c.lh rd' , uimm($rs1'$)	Load signed halfword, 16-bit encoding
✓	✓	c.sb $rs2'$, uimm($rs1'$)	Store byte, 16-bit encoding
✓	✓	c.sh $rs2'$, uimm($rs1'$)	Store halfword, 16-bit encoding

Zcyp

Zcyp is the set of sequenced instructions for code-size reduction.

This extension reuses some encodings from c.fld, c.fldsp, c.fsd, c.fsdsp. Therefore it is *incompatible* with D. It is fully compatible with F and, with Zdinx and also with [Zca](#).

NOTE

The PUSH/POP instructions share assembly mnemonics for different encodings. For further information see [PUSH/POP Register Instructions](#).

The PUSH/POP assembly syntax uses several variables, the meaning of which are:

- *reg_list* is a list containing 1 to 13 registers (ra and 0 to 12 s registers)
 - valid values: {ra}, {ra, s0}, {ra, s0-s1}, {ra, s0-s2}, ..., {ra, s0-s8}, {ra, s0-s9}, {ra, s0-s11}
 - note that {ra, s0-s10} is *not* valid, giving 12 lists not 13 for better encoding
- *stack_adj* is the total size of the stack frame.
 - valid values vary with register list length and the specific encoding, see the instruction pages for details.

RV32	RV64	Mnemonic	Instruction
✓	✓	c.push {reg_list}, -stack_adj	c.push : Create stack frame: push registers, allocate additional stack space.
✓	✓	c.pop {reg_list}, stack_adj	c.pop : Destroy stack frame: pop registers, deallocate stack frame.
✓	✓	c.popret {reg_list}, stack_adj	c.popret : Destroy stack frame: pop registers, deallocate stack frame, return.
✓	✓	c.popretz {reg_list}, stack_adj	c.popretz : Destroy stack frame: pop registers, deallocate stack frame, return zero.
✓	✓	c.mva01s sreg1, sreg2	c.mva01s : move two s0-s7 registers into a0-a1
✓	✓	c.mvsa01 sreg1, sreg2	c.mvsa01 : move a0-a1 into two different s0-s7 registers

Zcype

This extension reuses some encodings from c.fld, c.fldsp, c.fsd, c.fsdsp. Therefore it is *incompatible* with D.

Zcype offers EABI support for register mappings from [Zcyp](#) where the x register mapping is different to the UABI. The EABI specification is not frozen so these instructions cannot yet be accurately specified.

Zcyt

Zcyt is the set of table jump instructions for code-size reduction.

This extension reuses some encodings from c.fld, c.fldsp, c.fsd, c.fsdsp. Therefore it is *incompatible* with D. It is fully compatible with F and, with Zdinx and also with [Zca](#).

RV32	RV64	Mnemonic	Instruction
✓	✓	c.jt #index	c.jt: jump via table without link
✓	✓	c.jalt #index	c.jalt: jump via table and link to ra

c.zext.b

Synopsis

Zero extend byte, 16-bit encoding

Mnemonic

c.zext.b *rsd'*

Encoding (RV32, RV64)

```
Failed to generate image: wavedrom-cli.CMD failed:
{reg:[
  { bits: 2, name: 0x1, attr: ['OP=C1'] },
  { bits: 3, name: 0x0, attr: ['C.ZEXT.B'] },
  { bits: 2, name: 0x3, attr: ['FUNCT2'] },
  { bits: 3, name: 'rsd\\'', attr: ['SRCDEST'] },
  { bits: 6, name: 0x27, attr: ['FUNCT6'] },
],config:{bits:16}}
```

Description

This instruction takes a single source/destination operand, from the 8-register set x8-x15. It zero-extends the least-significant byte of the operand to XLEN by inserting zeros into all of the bits more significant than 7.

Prerequisites

C or Zca must also be configured.

32-bit equivalent

```
andi rsd, rsd, 0xff
```

NOTE

The SAIL module variable for *rsd'* is called *rsdc*.

Operation

```
X(rsdc) = EXTZ(X(rsdc)[7..0]);
```

Included in

Extension	Minimum version	Lifecycle state
Zcb (Zcb)	0.61.0	Stable

c.sext.b


Synopsis

Sign extend byte, 16-bit encoding

Mnemonic

c.sext.b *rsd'*

Encoding (RV32, RV64)



Description

This instruction takes a single source/destination operand, from the 8-register set x8-x15. It sign-extends the least-significant byte in the operand to XLEN by copying the most-significant bit in the byte (i.e., bit 7) to all of the more-significant bits.

Prerequisites

C or Zca must also be configured.

32-bit equivalent

[\[insns-sext_b\]](#) from Zbb

NOTE

The SAIL module variable for *rsd'* is called *rsdc*.

Operation

```
X(rsdc) = EXTS(X(rsdc)[7..0]);
```

Included in

Extension	Minimum version	Lifecycle state
Zcb (Zcb)	0.61.0	Stable

c.zext.h


Synopsis

Zero extend halfword, 16-bit encoding

Mnemonic

c.zext.h *rsd'*

Encoding (RV32, RV64)



Description

This instruction takes a single source/destination operand, from the 8-register set x8-x15. It zero-extends the least-significant halfword of the operand to XLEN by inserting zeros into all of the bits more significant than 15.

Prerequisites

C or Zca must also be configured.

32-bit equivalent

[\[insns-zext_h\]](#) from Zbb

NOTE

The SAIL module variable for *rsd'* is called *rsdc*.

Operation

```
X(rsdc) = EXTZ(X(rsdc)[15..0]);
```

Included in

Extension	Minimum version	Lifecycle state
Zcb (Zcb)	0.61.0	Stable

c.sext.h


Synopsis

Sign extend halfword, 16-bit encoding

Mnemonic

c.sext.h *rsd'*

Encoding (RV32, RV64)



Description

This instruction takes a single source/destination operand, from the 8-register set x8-x15. It sign-extends the least-significant halfword in the operand to XLEN by copying the most-significant bit in the halfword (i.e., bit 15) to all of the more-significant bits.

Prerequisites

C or Zca must also be configured.

32-bit equivalent

[\[insns-sext_h\]](#) from Zbb

NOTE

The SAIL module variable for *rsd'* is called *rsdc*.

Operation

```
X(rsdc) = EXTS(X(rsdc)[15..0]);
```

Included in

Extension	Minimum version	Lifecycle state
Zcb (Zcb)	0.61.0	Stable

c.zext.w

Synopsis

Zero extend word, 16-bit encoding

Mnemonic

c.zext.w *rsd'*

Encoding (RV64)

```
Failed to generate image: wavedrom-cli.CMD failed:
{reg:[
  { bits: 2, name: 0x1, attr: ['OP=C1'] },
  { bits: 3, name: 0x4, attr: ['C.ZEXT.W'] },
  { bits: 2, name: 0x3, attr: ['FUNCT2'] },
  { bits: 3, name: 'rsd\\'', attr: ['SRCDST'] },
  { bits: 6, name: 0x27, attr: ['FUNCT6'] },
],config:{bits:16}}
```

Description

This instruction takes a single source/destination operand, from the 8-register set x8-x15. It zero-extends the least-significant word of the operand to XLEN by inserting zeros into all of the bits more significant than 31.

Prerequisites

C or Zca must also be configured.

32-bit equivalent

```
add.uw rsd', rsd', zero
```

NOTE The SAIL module variable for *rsd'* is called *rsdc*.

Operation

```
X(rsdc) = EXTZ(X(rsdc)[31..0]);
```

Included in

Extension	Minimum version	Lifecycle state
Zcb (Zcb)	0.61.0	Stable

c.not


Synopsis

Bitwise not, 16-bit encoding

Mnemonic

c.not *rsd'*

Encoding (RV32, RV64)



Description

This instruction takes the one's complement of *rsd'* and writes the result to the same register. The operand is from the 8-register set x8-x15.

Prerequisites

C or Zca must also be configured.

32-bit equivalent

```
xori rd, rs, -1
```

NOTE

The SAIL module variable for *rsd'* is called *rsdc*.

Operation

$$X(rsdc) = X(rsdc) \text{ XOR } -1;$$

Included in

Extension	Minimum version	Lifecycle state
Zcb (Zcb)	0.61.0	Stable

c.mul


Synopsis

Multiply, 16-bit encoding

Mnemonic

c.mul *rsd'*, *rs2'*

Encoding (RV32, RV64)



Description

This instruction multiplies XLEN bits of the source operands from *rsd'* and *rs2'* and writes the lowest XLEN bits of the result to *rsd'*. Both operands are from the 8-register set x8-x15.

Prerequisites

C or Zca and either M or Zmmul must also be configured.

32-bit equivalent

[\[insns-mul\]](#)

NOTE

The SAIL module variable for *rsd'* is called *rsdc*, and for *rs2'* is called *rs2c*.

Operation

```
let result_wide = to_bits(2 * sizeof(xlen), signed(X(rsdc)) * signed(X(rs2c)));
X(rsdc) = result_wide[(sizeof(xlen) - 1) .. 0];
```

Included in

Extension	Minimum version	Lifecycle state
Zcb (Zcb)	0.61.0	Stable

c.lbu


Synopsis

Load unsigned byte, 16-bit encoding

Mnemonic

c.lbu *rd'*, *uimm(rs1')*

Encoding (RV32, RV64)



The immediate offset is formed as follows:

```
uimm[31:4] = 0;
uimm[3]     = encoding[10];
uimm[2:1]   = encoding[6:5];
uimm[0]     = encoding[11];
```

Description

This instruction loads a byte from the memory address formed by adding *rs1'* to the zero extended immediate *uimm*. The resulting byte is zero extended to XLEN bits and is written to *rd'*.

NOTE

rd' and *rs1'* are from the standard 8-register set x8-x15.

Prerequisites

Zca must also be configured, and not C to avoid an encoding conflict.

32-bit equivalent

[\[insns-lbu\]](#)

Operation

```
//This is not SAIL, it's pseudo-code. The SAIL hasn't been written yet.

X(rdc) = EXTZ(mem[X(rs1c)+EXTZ(uimm)] [7..0]);
```

Included in

Extension	Minimum version	Lifecycle state
Zcyh (Zcyh)	0.61.0	Stable

c.lhu

Synopsis

Load unsigned halfword, 16-bit encoding

Mnemonic

c.lhu *rd'*, *uimm(rs1')*

Encoding (RV32, RV64)

```
Failed to generate image: wavedrom-cli.CMD failed:
{reg:[
  { bits: 2, name: 0x2, attr: ['C0'] },
  { bits: 3, name: 'rd\' ' },
  { bits: 2, name: 'uimm[2:1]' },
  { bits: 3, name: 'rs1\' ' },
  { bits: 2, name: 'uimm[4:3]' },
  { bits: 1, name: 0x1 },
  { bits: 3, name: 0x1, attr: ['FUNCT3'] },
],config:{bits:16}}
```

The immediate offset is formed as follows:

```
uimm[31:5] = 0;
uimm[4:3]  = encoding[11:10];
uimm[2:1]  = encoding[6:5];
uimm[0]    = 0;
```

Description

This instruction loads a halfword from the memory address formed by adding *rs1'* to the zero extended immediate *uimm*. The resulting halfword is zero extended to XLEN bits and is written to *rd'*.

NOTE *rd'* and *rs1'* are from the standard 8-register set x8-x15.

Prerequisites

Zca must also be configured, and not C to avoid an encoding conflict.

32-bit equivalent

[\[insns-lhu\]](#)

Operation

```
//This is not SAIL, it's pseudo-code. The SAIL hasn't been written yet.

X(rdc) = EXTZ(load_mem[X(rs1c)+EXTZ(uimm)][15..0]);
```

Included in

Extension	Minimum version	Lifecycle state
Zcyh (Zcyh)	0.61.0	Stable

c.lb

Synopsis

Load signed byte, 16-bit encoding

Mnemonic

c.lb *rd'*, *uimm(rs1')*

Encoding (RV32, RV64)

15	13	12	11	10	9	7	6	5	4	2	1	0
0	0	1	0	uimm[0:3]		rs1'		uimm[2:1]		rd'	0	0
FUNCT3											C2	

The immediate offset is formed as follows:

```

uimm[31:4] = 0;
uimm[3]    = encoding[10];
uimm[2:1]  = encoding[6:5];
uimm[0]    = encoding[11];

```

Description

This instruction loads a byte from the memory address formed by adding *rs1'* to the zero extended immediate *uimm*. The resulting byte is sign extended to XLEN bits and is written to *rd'*.

NOTE | *rd'* and *rs1'* are from the standard 8-register set x8-x15.

Prerequisites

Zca must also be configured, and not C to avoid an encoding conflict.

32-bit equivalent

[\[insns-lb\]](#)

Operation

```
//This is not SAIL, it's pseudo-code. The SAIL hasn't been written yet.
```

```
X(rdc) = EXTS(mem[X(rs1c)+EXTZ(uimm)] [7..0]);
```

Included in

Extension	Minimum version	Lifecycle state
Zcyh (Zcyh)	0.61.0	Stable

c.lh

Synopsis

Load signed halfword, 16-bit encoding

Mnemonic

`c.lh rd', uimm(rs1')`

Encoding (RV32, RV64)

```
Failed to generate image: wavedrom-cli.CMD failed:
{reg:[
  { bits: 2, name: 0x0, attr: ['C0'] },
  { bits: 3, name: 'rd\'' },
  { bits: 2, name: 'uimm[2:1]' },
  { bits: 3, name: 'rs1\'' },
  { bits: 2, name: 'uimm[4:3]' },
  { bits: 1, name: 0x1 },
  { bits: 3, name: 0x1, attr: ['FUNCT3'] },
],config:{bits:16}}
```

The immediate offset is formed as follows:

```
uimm[31:5] = 0;
uimm[4:3]  = encoding[11:10];
uimm[2:1]  = encoding[6:5];
uimm[0]    = 0;
```

Description

This instruction loads a halfword from the memory address formed by adding *rs1'* to the zero extended immediate *uimm*. The resulting halfword is sign extended to XLEN bits and is written to *rd'*.

NOTE | *rd'* and *rs1'* are from the standard 8-register set x8-x15.

Prerequisites

Zca must also be configured, and not C to avoid an encoding conflict.

32-bit equivalent

[\[insns-lh\]](#)

Operation

```
//This is not SAIL, it's pseudo-code. The SAIL hasn't been written yet.

X(rdc) = EXTS(load_mem[X(rs1c)+EXTZ(uimm)][15..0]);
```

Included in

Extension	Minimum version	Lifecycle state
Zcyh (Zcyh)	0.61.0	Stable

c.sb


Synopsis

Store byte, 16-bit encoding

Mnemonic

c.sb *r2'*, *uimm(rs1')*

Encoding (RV32, RV64)



The immediate offset is formed as follows:

```
uimm[31:4] = 0;
uimm[3]     = encoding[10];
uimm[2:1]   = encoding[6:5];
uimm[0]     = encoding[11];
```

Description

This instruction stores the least significant byte of *rs2'* to the memory address formed by adding *rs1'* to the zero extended immediate *uimm*.

NOTE

rd' and *rs1'* are from the standard 8-register set x8-x15.

Prerequisites

Zca must also be configured, and not C to avoid an encoding conflict.

32-bit equivalent

[\[insns-sb\]](#)

Operation

```
//This is not SAIL, it's pseudo-code. The SAIL hasn't been written yet.

mem[X(rs1c)+EXTZ(uimm)] [7..0] = X(rs2c)
```

Included in

Extension	Minimum version	Lifecycle state
Zcyh (Zcyh)	0.61.0	Stable

c.sh

Synopsis

Store halfword, 16-bit encoding

Mnemonic

`c.sh $rs2'$, uimm($rs1'$)`

Encoding (RV32, RV64)

15	13	12	11	10	9	7	6	5	4	2	1	0
1	0	1	1	uimm[4:3]		$rs1'$		uimm[2:1]		$rs2'$	0	0
FUNCT3											C0	

The immediate offset is formed as follows:

```
uimm[31:5] = 0;
uimm[4:3]  = encoding[11:10];
uimm[2:1]  = encoding[6:5];
uimm[0]    = 0;
```

Description

This instruction stores the least significant halfword of $rs2'$ to the memory address formed by adding $rs1'$ to the zero extended immediate *uimm*.

NOTE rd' and $rs1'$ are from the standard 8-register set x8-x15.

Prerequisites

Zca must also be configured, and not C to avoid an encoding conflict.

32-bit equivalent

[\[insns-sh\]](#)

Operation

//This is not SAIL, it's pseudo-code. The SAIL hasn't been written yet.

```
mem[X(rs1c)+EXTZ(uimm)][15..0] = X(rs2c)
```

Included in

Extension	Minimum version	Lifecycle state
Zcyh (Zcyh)	0.61.0	Stable

c.push

Synopsis

Create stack frame: store ra and 0 to 12 saved registers to the stack frame, optionally allocate additional stack space.

Mnemonic

c.push

Encoding (RV32, RV64)

15	13	12		8	7		4	3	2	1	0
1	0	1	1	0	1	1	0	rlist	spimm[5:4]	1	0

NOTE rlist values 0 to 3 are reserved for a future EABI variant called *c.push.e*

Assembly Syntax

```
c.push {reg_list}, -stack_adj
c.push {xreg_list}, -stack_adj
```

The variables used in the assembly syntax are defined below.

```
RV32E:

switch (rlist){
  case 4: {reg_list="ra";      xreg_list="x1";}
  case 5: {reg_list="ra, s0";  xreg_list="x1, x8";}
  case 6: {reg_list="ra, s0-s1"; xreg_list="x1, x8-x9";}
  default: take_illegal_instruction_exception();
}

stack_adj      = stack_adj_base + spimm[5:4] * 16;
```

RV32I, RV64:

```
switch (rlist){
  case 4: {reg_list="ra";          xreg_list="x1";}
  case 5: {reg_list="ra, s0";      xreg_list="x1, x8";}
  case 6: {reg_list="ra, s0-s1";  xreg_list="x1, x8-x9";}
  case 7: {reg_list="ra, s0-s2";  xreg_list="x1, x8-x9, x18";}
  case 8: {reg_list="ra, s0-s3";  xreg_list="x1, x8-x9, x18-x19";}
  case 9: {reg_list="ra, s0-s4";  xreg_list="x1, x8-x9, x18-x20";}
  case 10: {reg_list="ra, s0-s5"; xreg_list="x1, x8-x9, x18-x21";}
  case 11: {reg_list="ra, s0-s6"; xreg_list="x1, x8-x9, x18-x22";}
  case 12: {reg_list="ra, s0-s7"; xreg_list="x1, x8-x9, x18-x23";}
  case 13: {reg_list="ra, s0-s8"; xreg_list="x1, x8-x9, x18-x24";}
  case 14: {reg_list="ra, s0-s9"; xreg_list="x1, x8-x9, x18-x25";}
  //note - to include s10, s11 must also be included
  case 15: {reg_list="ra, s0-s11"; xreg_list="x1, x8-x9, x18-x27";}
  default: take_illegal_instruction_exception();
}
stack_adj      = stack_adj_base + spimm[5:4] * 16;
```

RV32E:

```
stack_adj_base = 16;
Valid values:
stack_adj      = [16|32|48|64];
```

RV32I:

```
switch (rlist) {
  case 4.. 7: stack_adj_base = 16;
  case 8..11: stack_adj_base = 32;
  case 12..14: stack_adj_base = 48;
  case      15: stack_adj_base = 64;
}

Valid values:
switch (rlist) {
  case 4.. 7: stack_adj = [16|32|48| 64];
  case 8..11: stack_adj = [32|48|64| 80];
  case 12..14: stack_adj = [48|64|80| 96];
  case      15: stack_adj = [64|80|96|112];
}
```

RV64:

```
switch (rlist) {  
    case 4.. 5: stack_adj_base = 16;  
    case 6.. 7: stack_adj_base = 32;  
    case 8.. 9: stack_adj_base = 48;  
    case 10..11: stack_adj_base = 64;  
    case 12..13: stack_adj_base = 80;  
    case     14: stack_adj_base = 96;  
    case     15: stack_adj_base = 112;  
}
```

Valid values:

```
switch (rlist) {  
    case 4.. 5: stack_adj = [ 16| 32| 48| 64];  
    case 6.. 7: stack_adj = [ 32| 48| 64| 80];  
    case 8.. 9: stack_adj = [ 48| 64| 80| 96];  
    case 10..11: stack_adj = [ 64| 80| 96|112];  
    case 12..13: stack_adj = [ 80| 96|112|128];  
    case     14: stack_adj = [ 96|112|128|144];  
    case     15: stack_adj = [112|128|144|160];  
}
```


Description

This instruction pushes (stores) the registers in *reg_list* to the memory below the stack pointer, and then creates the stack frame by decrementing the stack pointer by *stack_adj*, including any additional stack space requested by the value of *spimm*.

NOTE

All ABI register mappings are for the UABI. An EABI version is planned once the EABI is frozen.

For further information see [PUSH/POP Register Instructions](#).

Stack Adjustment Calculation

stack_adj_base is the minimum number of bytes, in multiples of 16-byte blocks, required to cover the registers in the list.

spimm is the number of additional 16-byte blocks allocated for the stack frame.

The total stack adjustment represents the total size of the stack frame, which is *stack_adj_base* added to *spimm* scaled by 16, as defined above.

Prerequisites

C or Zca must also be configured.

32-bit equivalent

No direct equivalent encoding exists

Operation

The first section of pseudo-code may be executed multiple times before the instruction successfully completes.

```
//This is not SAIL, it's pseudo-code. The SAIL hasn't been written yet.

if (misa.MXL==1) bytes=4; else bytes=8;

addr=sp-bytes;
for(i in 27,26,25,24,23,22,21,20,19,18,9,8,1) {
    //if register i is in xreg_list
    if (xreg_list[i]) {
        switch(bytes) {
            4:  asm("sw x[i], 0(addr)");
            8:  asm("sd x[i], 0(addr)");
        }
        addr-=bytes;
    }
}
```

The final section of pseudo-code executes atomically, and only executes if the section above completes without any exceptions or interrupts.

```
//This is not SAIL, it's pseudo-code. The SAIL hasn't been written yet.

sp-=stack_adj;
```

RV32I Assembly example

```
c.push {ra, s0-s2}, -64
```

Encoding: *rlist*=7, *spimm*=3

The equivalent interrupt-safe instruction sequence is:

```
addi sp, sp, -64;
sw  s2, 60(sp);
sw  s1, 56(sp);
sw  s0, 52(sp);
sw  ra, 48(sp);
```

RV32I Assembly example

```
c.push {ra, s0-s1}, -32
```

Encoding: *rlist*=6, *spimm*=1

The equivalent interrupt-safe instruction sequence is:

```
addi sp, sp, -32;
sw  s1, 28(sp);
sw  s0, 24(sp);
sw  ra, 20(sp);
```

RV32I Assembly example

```
c.push {ra, s0-s3}, -64
```

Encoding: *rlist*=8, *spimm*=2

The equivalent interrupt-safe instruction sequence is:

```
addi sp, sp, -64;
sw  s3, 60(sp);
sw  s2, 56(sp);
sw  s1, 52(sp);
sw  s0, 48(sp);
sw  ra, 44(sp);
```

RV32I Assembly example

```
c.push {ra, s0-s11}, -112
```

Encoding: *rlist*=15, *spimm*=3

The equivalent interrupt-safe instruction sequence is:

```
addi sp, sp, -112;
sw  s11, 108(sp);
sw  s10, 104(sp);
sw  s9, 100(sp);
sw  s8, 96(sp);
sw  s7, 92(sp);
sw  s6, 88(sp);
sw  s5, 84(sp);
sw  s4, 80(sp);
sw  s3, 76(sp);
sw  s2, 72(sp);
sw  s1, 68(sp);
sw  s0, 64(sp);
sw  ra, 60(sp);
```

Included in

Extension	Minimum version	Lifecycle state
Zcyp (Zcyp)	0.61.0	Stable

Included in

Extension	Minimum version	Lifecycle state
Zcype (Zcype)	0.61.0	Stable

c.pop

Synopsis

Destroy stack frame: load ra and 0 to 12 saved registers from the stack frame, deallocate the stack frame.

Mnemonic

c.pop

Encoding (RV32, RV64)

15	13	12	8	7	4	3	2	1	0		
1	0	1	1	1	0	1	0	rlist	spimm[5:4]	1	0

NOTE

rlist values 0 to 3 are reserved for a future EABI variant called *c.pop.e*

Assembly Syntax

```
c.pop {reg_list}, stack_adj
c.pop {xreg_list}, stack_adj
```

The variables used in the assembly syntax are defined below.

RV32E:

```
switch (rlist){
    case 4: {reg_list="ra";      xreg_list="x1";}
    case 5: {reg_list="ra, s0";  xreg_list="x1, x8";}
    case 6: {reg_list="ra, s0-s1"; xreg_list="x1, x8-x9";}
    default: take_illegal_instruction_exception();
}
stack_adj      = stack_adj_base + spimm[5:4] * 16;
```

RV32I, RV64:

```
switch (rlist){
  case 4: {reg_list="ra";          xreg_list="x1";}
  case 5: {reg_list="ra, s0";      xreg_list="x1, x8";}
  case 6: {reg_list="ra, s0-s1";  xreg_list="x1, x8-x9";}
  case 7: {reg_list="ra, s0-s2";  xreg_list="x1, x8-x9, x18";}
  case 8: {reg_list="ra, s0-s3";  xreg_list="x1, x8-x9, x18-x19";}
  case 9: {reg_list="ra, s0-s4";  xreg_list="x1, x8-x9, x18-x20";}
  case 10: {reg_list="ra, s0-s5"; xreg_list="x1, x8-x9, x18-x21";}
  case 11: {reg_list="ra, s0-s6"; xreg_list="x1, x8-x9, x18-x22";}
  case 12: {reg_list="ra, s0-s7"; xreg_list="x1, x8-x9, x18-x23";}
  case 13: {reg_list="ra, s0-s8"; xreg_list="x1, x8-x9, x18-x24";}
  case 14: {reg_list="ra, s0-s9"; xreg_list="x1, x8-x9, x18-x25";}
  //note - to include s10, s11 must also be included
  case 15: {reg_list="ra, s0-s11"; xreg_list="x1, x8-x9, x18-x27";}
  default: take_illegal_instruction_exception();
}
stack_adj      = stack_adj_base + spimm[5:4] * 16;
```

RV32E:

```
stack_adj_base = 16;
Valid values:
stack_adj      = [16|32|48|64];
```

RV32I:

```
switch (rlist) {
  case 4.. 7: stack_adj_base = 16;
  case 8..11: stack_adj_base = 32;
  case 12..14: stack_adj_base = 48;
  case     15: stack_adj_base = 64;
}

Valid values:
switch (rlist) {
  case 4.. 7: stack_adj = [16|32|48| 64];
  case 8..11: stack_adj = [32|48|64| 80];
  case 12..14: stack_adj = [48|64|80| 96];
  case     15: stack_adj = [64|80|96|112];
}
```

RV64:

```
switch (rlist) {
  case 4.. 5: stack_adj_base = 16;
  case 6.. 7: stack_adj_base = 32;
  case 8.. 9: stack_adj_base = 48;
  case 10..11: stack_adj_base = 64;
  case 12..13: stack_adj_base = 80;
  case 14: stack_adj_base = 96;
  case 15: stack_adj_base = 112;
}
```

Valid values:

```
switch (rlist) {
  case 4.. 5: stack_adj = [ 16| 32| 48| 64];
  case 6.. 7: stack_adj = [ 32| 48| 64| 80];
  case 8.. 9: stack_adj = [ 48| 64| 80| 96];
  case 10..11: stack_adj = [ 64| 80| 96|112];
  case 12..13: stack_adj = [ 80| 96|112|128];
  case 14: stack_adj = [ 96|112|128|144];
  case 15: stack_adj = [112|128|144|160];
}
```

Description

This instruction pops (loads) the registers in *reg_list* from stack memory, and then adjusts the stack pointer by *stack_adj*.

NOTE

All ABI register mappings are for the UABI. An EABI version is planned once the EABI is frozen.

For further information see [PUSH/POP Register Instructions](#).

Stack Adjustment Calculation

stack_adj_base is the minimum number of bytes, in multiples of 16-byte blocks, required to cover the registers in the list.

spimm is the number of additional 16-byte blocks allocated for the stack frame.

The total stack adjustment represents the total size of the stack frame, which is *stack_adj_base* added to *spimm* scaled by 16, as defined above.

Prerequisites

C or Zca must also be configured.

32-bit equivalent

No direct equivalent encoding exists

Operation

The first section of pseudo-code may be executed multiple times before the instruction successfully completes.

```
//This is not SAIL, it's pseudo-code. The SAIL hasn't been written yet.

if (misa.MXL==1) bytes=4; else bytes=8;

addr=sp+stack_adj-bytes;
for(i in 27,26,25,24,23,22,21,20,19,18,9,8,1) {
    //if register i is in xreg_list
    if (xreg_list[i]) {
        switch(bytes) {
            4:  asm("lw x[i], 0(addr)");
            8:  asm("ld x[i], 0(addr)");
        }
        addr-=bytes;
    }
}
```

The final section of pseudo-code executes atomically, and only executes if the section above completes without any exceptions or interrupts.

```
//This is not SAIL, it's pseudo-code. The SAIL hasn't been written yet.

sp+=stack_adj;
```

RV32I Assembly example

```
c.pop    {ra}, 16
```

Encoding: *rlist*=4, *spimm*=0

The equivalent interrupt-safe instruction sequence is:

```
lw    ra, 12(sp);  
addi sp, sp, 16;
```

RV32I Assembly example

```
c.pop    {ra, s0-s2}, 48
```

Encoding: *rlist*=7, *spimm*=2

The equivalent interrupt-safe instruction sequence is:

```
lw    s2, 44(sp);  
lw    s1, 40(sp);  
lw    s0, 36(sp);  
lw    ra, 32(sp);  
addi sp, sp, 48;
```

RV32I Assembly example

```
c.pop {ra, s0-s3}, 48
```

Encoding: *rlist*=8, *spimm*=1

The equivalent interrupt-safe instruction sequence is:

```
lw    s3, 44(sp);  
lw    s2, 40(sp);  
lw    s1, 36(sp);  
lw    s0, 32(sp);  
lw    ra, 28(sp);  
addi sp, sp, 48;
```

RV32I Assembly example

```
c.pop {ra, s0-s4}, 64
```

Encoding: *rlist*=9, *spimm*=2

The equivalent interrupt-safe instruction sequence is:

```
lw    s4, 60(sp);  
lw    s3, 56(sp);  
lw    s2, 52(sp);  
lw    s1, 48(sp);  
lw    s0, 44(sp);  
lw    ra, 40(sp);  
addi  sp, sp, 64;
```

Included in

Extension	Minimum version	Lifecycle state
Zcyp (Zcyp)	0.61.0	Stable

Included in

Extension	Minimum version	Lifecycle state
Zcype (Zcype)	0.61.0	Stable

c.popret

Synopsis

Destroy stack frame: load ra and 0 to 12 saved registers from the stack frame, deallocate the stack frame, return to ra.

Mnemonic

c.popret

Encoding (RV32, RV64)

```
Failed to generate image: wavedrom-cli.CMD failed:
{reg:[
  { bits: 2, name: 0x2,          attr: [] },
  { bits: 2, name: 'spimm\[5:4\]', attr: [] },
  { bits: 4, name: 'rlist',      attr: [] },
  { bits: 5, name: 0x1e,        attr: [] },
  { bits: 3, name: 0x5,         attr: [] },
],config:{bits:16}}
```

NOTE

rlist values 0 to 3 are reserved for a future EABI variant called *c.popret.e*

Assembly Syntax

```
c.popret {reg_list}, stack_adj
c.popret {xreg_list}, stack_adj
```

The variables used in the assembly syntax are defined below.

RV32E:

```
switch (rlist){
  case 4: {reg_list="ra";      xreg_list="x1";}
  case 5: {reg_list="ra, s0";  xreg_list="x1, x8";}
  case 6: {reg_list="ra, s0-s1"; xreg_list="x1, x8-x9";}
  default: take_illegal_instruction_exception();
}
stack_adj      = stack_adj_base + spimm[5:4] * 16;
```

RV32I, RV64:

```
switch (rlist){
  case 4: {reg_list="ra";          xreg_list="x1";}
  case 5: {reg_list="ra, s0";      xreg_list="x1, x8";}
  case 6: {reg_list="ra, s0-s1";   xreg_list="x1, x8-x9";}
  case 7: {reg_list="ra, s0-s2";   xreg_list="x1, x8-x9, x18";}
  case 8: {reg_list="ra, s0-s3";   xreg_list="x1, x8-x9, x18-x19";}
  case 9: {reg_list="ra, s0-s4";   xreg_list="x1, x8-x9, x18-x20";}
  case 10: {reg_list="ra, s0-s5";  xreg_list="x1, x8-x9, x18-x21";}
  case 11: {reg_list="ra, s0-s6";  xreg_list="x1, x8-x9, x18-x22";}
  case 12: {reg_list="ra, s0-s7";  xreg_list="x1, x8-x9, x18-x23";}
  case 13: {reg_list="ra, s0-s8";  xreg_list="x1, x8-x9, x18-x24";}
  case 14: {reg_list="ra, s0-s9";  xreg_list="x1, x8-x9, x18-x25";}
  //note - to include s10, s11 must also be included
  case 15: {reg_list="ra, s0-s11"; xreg_list="x1, x8-x9, x18-x27";}
  default: take_illegal_instruction_exception();
}
stack_adj      = stack_adj_base + spimm[5:4] * 16;
```

RV32E:

```
stack_adj_base = 16;
Valid values:
stack_adj      = [16|32|48|64];
```

RV32I:

```
switch (rlist) {
  case 4.. 7: stack_adj_base = 16;
  case 8..11: stack_adj_base = 32;
  case 12..14: stack_adj_base = 48;
  case      15: stack_adj_base = 64;
}

Valid values:
switch (rlist) {
  case 4.. 7: stack_adj = [16|32|48| 64];
  case 8..11: stack_adj = [32|48|64| 80];
  case 12..14: stack_adj = [48|64|80| 96];
  case      15: stack_adj = [64|80|96|112];
}
```

RV64:

```
switch (rlist) {  
    case 4.. 5: stack_adj_base = 16;  
    case 6.. 7: stack_adj_base = 32;  
    case 8.. 9: stack_adj_base = 48;  
    case 10..11: stack_adj_base = 64;  
    case 12..13: stack_adj_base = 80;  
    case     14: stack_adj_base = 96;  
    case     15: stack_adj_base = 112;  
}
```

Valid values:

```
switch (rlist) {  
    case 4.. 5: stack_adj = [ 16| 32| 48| 64];  
    case 6.. 7: stack_adj = [ 32| 48| 64| 80];  
    case 8.. 9: stack_adj = [ 48| 64| 80| 96];  
    case 10..11: stack_adj = [ 64| 80| 96|112];  
    case 12..13: stack_adj = [ 80| 96|112|128];  
    case     14: stack_adj = [ 96|112|128|144];  
    case     15: stack_adj = [112|128|144|160];  
}
```

Description

This instruction pops (loads) the registers in *reg_list* from stack memory, adjusts the stack pointer by *stack_adj* and then returns to *ra*.

NOTE

All ABI register mappings are for the UABI. An EABI version is planned once the EABI is frozen.

For further information see [PUSH/POP Register Instructions](#).

Stack Adjustment Calculation

stack_adj_base is the minimum number of bytes, in multiples of 16-byte blocks, required to cover the registers in the list.

spimm is the number of additional 16-byte blocks allocated for the stack frame.

The total stack adjustment represents the total size of the stack frame, which is *stack_adj_base* added to *spimm* scaled by 16, as defined above.

Prerequisites

C or Zca must also be configured.

32-bit equivalent

No direct equivalent encoding exists

Operation

The first section of pseudo-code may be executed multiple times before the instruction successfully completes.

```
//This is not SAIL, it's pseudo-code. The SAIL hasn't been written yet.

if (misa.MXL==1) bytes=4; else bytes=8;

addr=sp+stack_adj-bytes;
for(i in 27,26,25,24,23,22,21,20,19,18,9,8,1) {
    //if register i is in xreg_list
    if (xreg_list[i]) {
        switch(bytes) {
            4:  asm("lw x[i], 0(addr)");
            8:  asm("ld x[i], 0(addr)");
        }
        addr-=bytes;
    }
}
```

The final section of pseudo-code executes atomically, and only executes if the section above completes without any exceptions or interrupts.

```
//This is not SAIL, it's pseudo-code. The SAIL hasn't been written yet.

sp+=stack_adj;
asm("ret");
```


RV32I Assembly example

```
c.popret {ra}, 16
```

Encoding: *rlist*=4, *spimm*=0

The equivalent interrupt-safe instruction sequence is:

```
lw    ra, 12(sp);
addi  sp, sp, 16;
ret;
```

RV32I Assembly example

```
c.popret {ra, s0-s2}, 48
```

Encoding: *rlist*=7, *spimm*=2

The equivalent interrupt-safe instruction sequence is:

```
lw    s2, 44(sp);
lw    s1, 40(sp);
lw    s0, 36(sp);
lw    ra, 32(sp);
addi  sp, sp, 48;
ret;
```

RV32I Assembly example

```
c.popret {ra, s0-s3}, 48
```

Encoding: *rlist*=8, *spimm*=1

The equivalent interrupt-safe instruction sequence is:

```
lw    s3, 44(sp);
lw    s2, 40(sp);
lw    s1, 36(sp);
lw    s0, 32(sp);
lw    ra, 28(sp);
addi  sp, sp, 48;
ret;
```

RV32I Assembly example

```
c.popret {ra, s0-s4}, 64
```

Encoding: *rlist*=9, *spimm*=2

The equivalent interrupt-safe instruction sequence is:

```
lw    s4, 60(sp);
lw    s3, 56(sp);
lw    s2, 52(sp);
lw    s1, 48(sp);
lw    s0, 44(sp);
lw    ra, 40(sp);
addi  sp, sp, 64;
ret;
```

Included in

Extension	Minimum version	Lifecycle state
Zcyp (Zcyp)	0.61.0	Stable

Included in

Extension	Minimum version	Lifecycle state
Zcype (Zcype)	0.61.0	Stable

c.popretz

Synopsis

Destroy stack frame: load ra and 0 to 12 saved registers from the stack frame, deallocate the stack frame, move zero into a0, return to ra.

Mnemonic

c.popretz

Encoding (RV32, RV64)

15	13	12				8	7			4	3	2	1	0
1	0	1	1	1	1	0	0		rlist			spimm[5:4]	1	0

NOTE | *rlist* values 0 to 3 are reserved for a future EABI variant called *c.popretz.e*

Assembly Syntax

```
c.popretz {reg_list}, stack_adj
c.popretz {xreg_list}, stack_adj
```

```
RV32E:

switch (rlist){
    case 4: {reg_list="ra";      xreg_list="x1";}
    case 5: {reg_list="ra, s0";  xreg_list="x1, x8";}
    case 6: {reg_list="ra, s0-s1"; xreg_list="x1, x8-x9";}
    default: take_illegal_instruction_exception();
}
stack_adj      = stack_adj_base + spimm[5:4] * 16;
```

RV32I, RV64:

```
switch (rlist){
  case 4: {reg_list="ra";          xreg_list="x1";}
  case 5: {reg_list="ra, s0";      xreg_list="x1, x8";}
  case 6: {reg_list="ra, s0-s1";  xreg_list="x1, x8-x9";}
  case 7: {reg_list="ra, s0-s2";  xreg_list="x1, x8-x9, x18";}
  case 8: {reg_list="ra, s0-s3";  xreg_list="x1, x8-x9, x18-x19";}
  case 9: {reg_list="ra, s0-s4";  xreg_list="x1, x8-x9, x18-x20";}
  case 10: {reg_list="ra, s0-s5"; xreg_list="x1, x8-x9, x18-x21";}
  case 11: {reg_list="ra, s0-s6"; xreg_list="x1, x8-x9, x18-x22";}
  case 12: {reg_list="ra, s0-s7"; xreg_list="x1, x8-x9, x18-x23";}
  case 13: {reg_list="ra, s0-s8"; xreg_list="x1, x8-x9, x18-x24";}
  case 14: {reg_list="ra, s0-s9"; xreg_list="x1, x8-x9, x18-x25";}
  //note - to include s10, s11 must also be included
  case 15: {reg_list="ra, s0-s11"; xreg_list="x1, x8-x9, x18-x27";}
  default: take_illegal_instruction_exception();
}
stack_adj      = stack_adj_base + spimm[5:4] * 16;
```

RV32E:

```
stack_adj_base = 16;
Valid values:
stack_adj      = [16|32|48|64];
```

RV32I:

```
switch (rlist) {
  case 4.. 7: stack_adj_base = 16;
  case 8..11: stack_adj_base = 32;
  case 12..14: stack_adj_base = 48;
  case      15: stack_adj_base = 64;
}

Valid values:
switch (rlist) {
  case 4.. 7: stack_adj = [16|32|48| 64];
  case 8..11: stack_adj = [32|48|64| 80];
  case 12..14: stack_adj = [48|64|80| 96];
  case      15: stack_adj = [64|80|96|112];
}
```

RV64:

```
switch (rlist) {  
    case 4.. 5: stack_adj_base = 16;  
    case 6.. 7: stack_adj_base = 32;  
    case 8.. 9: stack_adj_base = 48;  
    case 10..11: stack_adj_base = 64;  
    case 12..13: stack_adj_base = 80;  
    case 14: stack_adj_base = 96;  
    case 15: stack_adj_base = 112;  
}
```

Valid values:

```
switch (rlist) {  
    case 4.. 5: stack_adj = [ 16| 32| 48| 64];  
    case 6.. 7: stack_adj = [ 32| 48| 64| 80];  
    case 8.. 9: stack_adj = [ 48| 64| 80| 96];  
    case 10..11: stack_adj = [ 64| 80| 96|112];  
    case 12..13: stack_adj = [ 80| 96|112|128];  
    case 14: stack_adj = [ 96|112|128|144];  
    case 15: stack_adj = [112|128|144|160];  
}
```

Description

This instruction pops (loads) the registers in *reg_list* from stack memory, adjusts the stack pointer by *stack_adj*, moves zero into a0 and then returns to *ra*.

NOTE

All ABI register mappings are for the UABI. An EABI version is planned once the EABI is frozen.

For further information see [PUSH/POP Register Instructions](#).

Stack Adjustment Calculation

stack_adj_base is the minimum number of bytes, in multiples of 16-byte blocks, required to cover the registers in the list.

spimm is the number of additional 16-byte blocks allocated for the stack frame.

The total stack adjustment represents the total size of the stack frame, which is *stack_adj_base* added to *spimm* scaled by 16, as defined above.

Prerequisites

C or Zca must also be configured.

32-bit equivalent

No direct equivalent encoding exists

Operation

The first section of pseudo-code may be executed multiple times before the instruction successfully completes.

```
//This is not SAIL, it's pseudo-code. The SAIL hasn't been written yet.

if (misa.MXL==1) bytes=4; else bytes=8;

addr=sp+stack_adj-bytes;
for(i in 27,26,25,24,23,22,21,20,19,18,9,8,1) {
    //if register i is in xreg_list
    if (xreg_list[i]) {
        switch(bytes) {
            4:  asm("lw x[i], 0(addr)");
            8:  asm("ld x[i], 0(addr)");
        }
        addr-=bytes;
    }
}
```

The final section of pseudo-code executes atomically, and only executes if the section above completes without any exceptions or interrupts.

NOTE

The *li a0, 0* **could** be executed more than once, but is included in the atomic section for convenience.

```
//This is not SAIL, it's pseudo-code. The SAIL hasn't been written yet.

asm("li a0, 0");
sp+=stack_adj;
asm("ret");
```

RV32I Assembly example

```
c.popretz {ra}, 16
```

Encoding: *rlist=4, spimm=0*

The equivalent interrupt-safe instruction sequence is:

```
lw    ra, 12(sp);  
li    a0, 0;  
addi  sp, sp, 16;  
ret;
```

RV32I Assembly example

```
c.popretz {ra, s0-s2}, 48
```

Encoding: *rlist=7, spimm=2*

The equivalent interrupt-safe instruction sequence is:

```
lw    s2, 44(sp);  
lw    s1, 40(sp);  
lw    s0, 36(sp);  
lw    ra, 32(sp);  
li    a0, 0;  
addi  sp, sp, 48;  
ret;
```

RV32I Assembly example

```
c.popretz {ra, s0-s3}, 48
```

Encoding: *rlist=8, spimm=1*

The equivalent interrupt-safe instruction sequence is:

```
lw    s3, 44(sp);  
lw    s2, 40(sp);  
lw    s1, 36(sp);  
lw    s0, 32(sp);  
lw    ra, 28(sp);  
li    a0, 0;  
addi  sp, sp, 48;  
ret;
```


RV32I Assembly example

```
c.popretz {ra, s0-s4}, 64
```

Encoding: *rlist*=9, *spimm*=2

The equivalent interrupt-safe instruction sequence is:

```
lw    s4, 60(sp);
lw    s3, 56(sp);
lw    s2, 52(sp);
lw    s1, 48(sp);
lw    s0, 44(sp);
lw    ra, 40(sp);
li    a0, 0;
addi  sp, sp, 64;
ret;
```

Included in

Extension	Minimum version	Lifecycle state
Zcyp (Zcyp)	0.61.0	Stable

Included in

Extension	Minimum version	Lifecycle state
Zcype (Zcype)	0.61.0	Stable

c.mva01s

Synopsis

Move two s0-s7 registers into a0-a1

Mnemonic

c.mva01s

Encoding (RV32, RV64)

15	13	12	10	9	7	6	5	4	2	1	0
1	0	1	0	1	1	sreg1	1	1	sreg2	1	0

Assembly Syntax

```
c.mva01s sreg1, sreg2
```

Description

This instruction moves *sreg1* into *a0* and *sreg2* into *a1*. The execution is atomic, so it is not possible to observe state where only one of *a0* or *a1* have been updated.

The encoding has uses *sreg* number specifiers instead of *xreg* number specifiers to save encoding space. The mapping between them is specified in the pseudo-code below.

NOTE

The *s* register mapping is taken from the UABI, and may not match the currently unratified EABI. *c.mva01s.e* may be included in the future.

Prerequisites

C or Zca must also be configured.

32-bit equivalent

No direct equivalent encoding exists.

Operation

```
//This is not SAIL, it's pseudo-code. The SAIL hasn't been written yet.

if (RV32E && (sreg1>1 || sreg2>1)) {
    take_illegal_instruction_exception();
}

xreg1 = {sreg1[2:1]>0,sreg1[2:1]==0,sreg1[2:0]};
xreg2 = {sreg2[2:1]>0,sreg2[2:1]==0,sreg2[2:0]};

X[10] = X[xreg1];
X[11] = X[xreg2];
```

Included in

Extension	Minimum version	Lifecycle state
Zcyp (Zcyp)	0.61.0	Stable

c.mvsa01

Synopsis

Move two s0-s7 registers into a0-a1

Mnemonic

c.mvsa01

Encoding (RV32, RV64)

15	13	12	10	9	7	6	5	4	2	1	0		
1	0	1	0	1	1	sreg1		0	1	sreg2		1	0

NOTE

For the encoding to be legal *sreg1* != *sreg2*.

Assembly Syntax

```
c.mvsa01 sreg1, sreg2
```

Description

This instruction moves *a0* into *sreg1* and *a1* and *sreg2*. *sreg1* and *sreg2* must be different. The execution is atomic, so it is not possible to observe state where only one of *sreg1* or *sreg2* has been updated.

The encoding has uses *sreg* number specifiers instead of *xreg* number specifiers to save encoding space. The mapping between them is specified in the pseudo-code below.

NOTE

The *s* register mapping is taken from the UABI, and may not match the currently unratified EABI. *c.mvsa01.e* may be included in the future.

Prerequisites

C or Zca must also be configured.

32-bit equivalent

No direct equivalent encoding exists.

Operation

```
//This is not SAIL, it's pseudo-code. The SAIL hasn't been written yet.

if (RV32E && (sreg1>1 || sreg2>1)) {
    take_illegal_instruction_exception();
}

xreg1 = {sreg1[2:1]>0,sreg1[2:1]==0,sreg1[2:0]};
xreg2 = {sreg2[2:1]>0,sreg2[2:1]==0,sreg2[2:0]};

X[xreg1] = X[10];
X[xreg2] = X[11];
```

Included in

Extension	Minimum version	Lifecycle state
Zcyp (Zcyp)	0.61.0	Stable

PUSH/POP register instructions

These instructions are collectively referred to as PUSH/POP:

- `c.push`: Create stack frame: push registers, allocate additional stack space.
- `c.pop`: Destroy stack frame: pop registers, deallocate stack frame.
- `c.popret`: Destroy stack frame: pop registers, deallocate stack frame, return.
- `c.popretz`: Destroy stack frame: pop registers, deallocate stack frame, return zero.

The term PUSH refers to `c.push`.

The term POP refers to `c.pop`.

The term POPRET refers to `c.popret` and `c.popretz`.

Common details for these instructions are in this section.

PUSH/POP functional overview

PUSH, POP, POPRET are used to reduce the size of function prologues and epilogues.

1. The PUSH instruction
 - pushes (stores) the registers specified in the register list to the stack frame
 - adjusts the stack pointer to create the stack frame
2. The POP instruction
 - pops (loads) the registers in the register list from the stack frame
 - adjusts the stack pointer to destroy the stack frame
3. The POPRET instructions
 - pop (load) the registers in the register list from the stack from
 - `c.popretz` also moves zero into `a0` as the return value
 - adjust the stack pointer to destroy the stack frame
 - execute a `ret` instruction to return from the function

Example usage

This example gives an illustration of the use of PUSH and POPRET.

The function *processMarkers* in the EMBench benchmark picojpeg in the following file on github: [libpicojpeg.c](#)

The prologue and epilogue compile with GCC10 to:

```
0001098a <processMarkers>:
1098a:      711d                addi    sp,sp,-96 ;#c.push(1)
1098c:      c8ca                sw      s2,80(sp) ;#c.push(2)
1098e:      c6ce                sw      s3,76(sp) ;#c.push(3)
10990:      c4d2                sw      s4,72(sp) ;#c.push(4)
10992:      ce86                sw      ra,92(sp) ;#c.push(5)
10994:      cca2                sw      s0,88(sp) ;#c.push(6)
10996:      caa6                sw      s1,84(sp) ;#c.push(7)
10998:      c2d6                sw      s5,68(sp) ;#c.push(8)
1099a:      c0da                sw      s6,64(sp) ;#c.push(9)
1099c:      de5e                sw      s7,60(sp) ;#c.push(10)
1099e:      dc62                sw      s8,56(sp) ;#c.push(11)
109a0:      da66                sw      s9,52(sp) ;#c.push(12)
109a2:      d86a                sw      s10,48(sp) ;#c.push(13)
109a4:      d66e                sw      s11,44(sp) ;#c.push(14)
...
109f4:      4501                li      a0,0 ;#c.popretz(1)
109f6:      40f6                lw      ra,92(sp) ;#c.popretz(2)
109f8:      4466                lw      s0,88(sp) ;#c.popretz(3)
109fa:      44d6                lw      s1,84(sp) ;#c.popretz(4)
109fc:      4946                lw      s2,80(sp) ;#c.popretz(5)
109fe:      49b6                lw      s3,76(sp) ;#c.popretz(6)
10a00:      4a26                lw      s4,72(sp) ;#c.popretz(7)
10a02:      4a96                lw      s5,68(sp) ;#c.popretz(8)
10a04:      4b06                lw      s6,64(sp) ;#c.popretz(9)
10a06:      5bf2                lw      s7,60(sp) ;#c.popretz(10)
10a08:      5c62                lw      s8,56(sp) ;#c.popretz(11)
10a0a:      5cd2                lw      s9,52(sp) ;#c.popretz(12)
10a0c:      5d42                lw      s10,48(sp) ;#c.popretz(13)
10a0e:      5db2                lw      s11,44(sp) ;#c.popretz(14)
10a10:      6125                addi    sp,sp,96 ;#c.popretz(15)
10a12:      8082                ret                    ;#c.popretz(16)
```

with the GCC option `-msave-restore` the output is the following:

```
0001080e <processMarkers>:
    1080e:      73a012ef          jal    t0,11f48 <__riscv_save_12>
    10812:      1101             addi   sp,sp,-32
    ...
    10862:      4501             li     a0,0
    10864:      6105             addi   sp,sp,32
    10866:      71e0106f          j      11f84 <__riscv_restore_12>
```

with PUSH/POPRET this reduces to

```
0001080e <processMarkers>:
    1080e:      xxxx             c.push {ra,s0-s11},-96
    ...
    10866:      xxxx             c.popretz {ra,s0-s11}, 96
```

The prologue / epilogue reduce from 58-bytes in the original code, to 16-bytes with `-msave-restore`, and to 6-bytes with PUSH and POPRET. As well as reducing the code-size PUSH and POPRET eliminate the branches from calling the millicode *save/restore* routines and so also perform better.

NOTE

The calls to `<riscv_save_0>/<riscv_restore_0>` become 64-bit when the target functions are out of the $\pm 1\text{MB}$ range, increasing the prologue/epilogue size to 22-bytes.

NOTE

POP is typically used in tail-calling sequences where *ret* is not used to return to *ra* after destroying the stack frame.

Compiler implementation

The technique used in the initial implementation in LLVM is to let the compiler generate the function prologue and epilogue, and then replace the instruction sequences with the relevant PUSH/POP instructions.

spimm handling

The instructions have a restricted range of *spimm* available. If this is insufficient then a separate `c.addi16sp` can be used to increase the range.

register list handling

The instructions do not directly support `{ra, s0-s10}` to reduce the amount of encoding space required. If this register list is required then *s11* should also be included. This costs a small amount of memory and performance, but saves code-size.

PUSH/POP Fault handling

The sequence required to execute the PUSH/POP instruction may be interrupted, or may not be able to start execution for several reasons.

- virtual memory page fault or PMP fault
 - these can be detected before execution, or during execution if the memory addresses cross a page/PMP boundary
 - xTVAL is set to any address which causes the fault
- watchpoint trigger
 - these can be detected before execution, or during execution depending on the trigger type (load data triggers require the sequence to have started executing, for example)
 - xTVAL is set to any address which causes the fault
- external debug halt
 - the halt can treat the whole sequence atomically, or interrupt mid sequence (implementation defined)
- debug halt caused by a trigger
 - same comment as watchpoint trigger above
- load access fault
 - these are detected while the sequence is executing
 - xTVAL is set to the fault address.
- store access fault (precise or imprecise)
 - these may be detected while the sequence is executing, or afterwards if imprecise
 - xTVAL is set to the fault address.
- interrupts
 - these may arrive at any time. An implementation can choose whether to interrupt the sequence or not.

NOTE

xTVAL may be hardwired to zero in an implementation. Recovering from faults such as page faults requires that it is implemented.

In all cases MEPC contains the PC of the PUSH/POP instruction, and MCAUSE is set as expected for the type of fault.

For debug halts DPC is set to the PC of the PUSH/POP instruction.

Because some faults can only be detected during the sequence the core implementation must be able to recover from the fault and re-execute the sequence. This may involve executing some or all of the loads and stores from the sequence multiple times before the sequence completes (as multiple faults or multiple interrupts are possible).

Therefore correct execution requires that *sp* refers to idempotent memory (also see [Non-idempotent memory handling](#)).

Software view of execution

Software view of the PUSH sequence

From a software perspective the PUSH sequence appears as:

- A sequence of stores writing a contiguous block of memory. Any of the bytes may be written multiple times.
- A stack pointer adjustment

Because the memory is idempotent and the stores are non-overlapping, they may be reordered, grouped into larger accesses, split into smaller access or any combination of these.

If an implementation allows interrupts during the sequence, and the interrupt handler uses *sp* to allocate stack memory, then any stores which were executed before the interrupt may be overwritten by the handler. This is safe because the memory is idempotent and the stores will be re-executed when execution resumes.

The stack pointer adjustment must only be committed once it is certain that all of the stores will complete without triggering any precise faults (for example, page faults). Stores may also return imprecise faults from the bus. It is platform defined whether the core implementation waits for the bus responses before continuing to the final stage of the sequence, or handles errors responses after completing the PUSH instruction.

For example:

```
c.push {ra, s0-s5}, -64
```

Appears to software as:

```
# any bytes from sp-1 to sp-28 may be written multiple times before the
# instruction completes
# therefore these updates may be visible in the interrupt/exception handler below
# the stack pointer
sw s5, -4(sp)
sw s4, -8(sp)
sw s3, -12(sp)
sw s2, -16(sp)
sw s1, -20(sp)
sw s0, -24(sp)
sw ra, -28(sp)

# this must only execute once, and will only execute after all stores completed
# without any precise faults
# therefore this update is only visible in the interrupt/exception handler if
# c.push has completed
addi sp, sp, -64
```

Software view of the POP/POPRET sequence

From a software perspective the POP/POPRET sequence appears as:

- A sequence of loads, any of which may be executed multiple times
- A stack pointer adjustment
- An optional LI zero into a0
- An optional RET

If an implementation allows interrupts during the sequence, then any loads which were executed before the interrupt may update architectural state. The loads will be re-executed once the handler completes, so the values will be overwritten. Therefore it is permitted for an implementation to update some of the destination registers before taking the interrupt or other fault.

The optional load immediate and stack pointer adjustment must only be committed once it is certain that all of the loads will complete successfully.

For POPRET once the stack pointer adjustment has been committed the RET must execute.

For example:

```
c.popretz {ra, s0-s3}, 32;
```

Appears to software as:

```
# any or all of these load instructions may execute multiple times
# therefore these updates may be visible in the interrupt/exception handler
lw    s3, 28(sp)
lw    s2, 24(sp)
lw    s1, 20(sp)
lw    s0, 16(sp)
lw    ra, 12(sp)

# these must only execute once, will only execute after all loads complete
# successfully
# all instructions must execute atomically
# therefore these updates are not visible in the interrupt/exception handler
li a0, 0
addi sp, sp, 32
ret
```

Forward progress guarantee

The PUSH/POP sequence has the same forward progress guarantee as executing the instructions from the equivalent assembly sequences.

Non-idempotent memory handling

An implementation may have a requirement to issue a PUSH/POP instruction to non-idempotent memory.

Error detection

If the core implementation does not support PUSH/POP to non-idempotent memories, the core may use an idempotency PMA to detect it and take a load (POP/POPRET) or store (PUSH) access fault exception in order to avoid unpredictable results.

Non-idempotent support

It is possible to support non-idempotent memory. One reason is to re-use PUSH/POP as a restricted form of a load/store multiple instruction to a peripheral, as there is no generic load/store multiple instruction in the RISC-V ISA.

If accessing non-idempotent memory then it is *recommended* to:

- 1. Not allow interrupts during execution
- 2. Not allow external debug halt during execution
- 3. Detect any virtual memory page faults or PMP faults for the whole instruction before starting execution (instead of during the sequence)
- 4. Not split / merge / reorder the generated memory accesses

It is possible that one of the following will still occur during execution:

- 1. Watchpoint trigger
- 2. Load/store access fault

In these cases the core will jump to the debug or exception handler. If execution is required to continue afterwards (so the event is not fatal to the code execution), then the handler is required to do so in software.

By following these rules memory accesses will only ever be issued once, and decreasing address order.

It is possible for implementations to follow these restricted rules and to safely access both types of memory. It is also possible for an implementation to use PMAs to detect the memory type and apply different rules, such as only allowing interrupts if accessing cacheable memory, for example.

Included in

Extension	Minimum version	Lifecycle state
Zcyp (Zcyp)	0.61.0	Stable

c.jt

Synopsis

jump via table without link

Mnemonic

c.jt

Encoding (RV32, RV64)

15	13	12	10	9						2	1	0
1	0	1	0	0	1	index8					1	0

NOTE

For this encoding to decode as *c.jt*, *index8*<64, otherwise it decodes as *c.jalt*: jump via table and link to ra.

Assembly Syntax

```
c.jt #index
```

Description

This instruction is used to dereference a table of PCs, and then jumps without linking to the dereferenced PC.

For further information see [Table Jump Instructions](#).

Prerequisites

C or Zca must also be configured.

32-bit equivalent

No direct equivalent encoding exists.

Operation

```
//This is not SAIL, it's pseudo-code. The SAIL hasn't been written yet.

# target_address is temporary internal state, it doesn't represent a real register
# Mem is byte indexed
# "index8" is the field from the encoding, not "index" passed to the C.JT/C.JALT
  in the assembler
# which is formed below

if (OPCODE=="C.JALT") {
    index = index8 - 64;
} else {
    index = index8;
}

switch(XLEN) {
    32: table_address[XLEN-1:0] = JVT.base + (index<<2);
    64: table_address[XLEN-1:0] = JVT.base + (index<<3);
}

//check for debug mode entry, trigger with timing=0 and action=1, haltreq or step
if ((debug_trigger(table_address) && MCONTROL.timing==0 && MCONTROL.action==1) ||
    external_debug_haltreq() || DCSR.step==1) {
    DPC          = current_PC;
    DCSR.cause = DCSR.step==1 ? 4 : external_debug_haltreq() ? 3 : 2;
    enter_debug_mode();
}
//check for breakpoint trigger which takes an exception with timing=0
} else if ((debug_trigger(table_address) && MCONTROL.timing==0) ||
            !can_access_instruction_memory(table_address)) {
    MEPC  = current_PC;
    MTVAL = table_address;
    MCAUSE = debug_trigger(table_address) ? BREAKPOINT : INSTRUCTION_ACCESS_FAULT;
    take_exception();
} else {
    //access the jump table
    switch(XLEN) {
        32: LW target_address, InstMemory[table_address][XLEN-1:0];
        64: LD target_address, InstMemory[table_address][XLEN-1:0];
    }

    //don't use haltreq or step here, only check the addresses
    //check for table_address after reading if timing=1
    if (debug_trigger(table_address) && MCONTROL.timing==1 && MCONTROL.action==1) {
        DPC          = current_PC;
        DCSR.cause = 2;
        enter_debug_mode();
    } else if (debug_trigger(table_address) && MCONTROL.timing==1) {
```

```
MEPC      = current_PC;
MTVAL     = table_address;
MCAUSE    = BREAKPOINT;
take_exception();
} else if ((debug_trigger(target_address) && MCONTROL.timing==0 &&
MCONTROL.action==1) {
    DPC      = target_address;
    DCSR.cause = 2;
    enter_debug_mode();
} else if (((debug_trigger(target_address) && MCONTROL.timing==0) ||
            !can_access_instruction_memory(target_address)) {
    MEPC      = target_address;
    MTVAL     = target_address;
    MCAUSE    = debug_trigger(target_address) ? BREAKPOINT :
INSTRUCTION_ACCESS_FAULT;
    take_exception();
} else {
    //jump to the target address
    if (OPCODE=="C.JALT") {
        JALR ra, target_address[XLEN-1:0]&~0x1;
    } else {
        JR target_address[XLEN-1:0]&~0x1;
    }
}
}
```

Included in

Extension	Minimum version	Lifecycle state
Zcyt (Zcyt)	0.61.0	Stable

c.jalt

Synopsis

jump via table and link to ra

Mnemonic

c.jalt

Encoding (RV32, RV64)

15	13	12	10	9						2	1	0
1	0	1	0	0	1				index8		1	0

NOTE

For this encoding to decode as *c.jalt*, $index8 \geq 64$, otherwise it decodes as [c.jt: jump via table without link](#).

Assembly Syntax

```
c.jalt #index
```

NOTE

index in the assembly syntax is valid from 0-192. *index8* in the encoding is valid from 64-255, so $index = index8 - 64$.

Description

This instruction is used to dereference a table of PCs, and then jumps to the dereferenced PC and links to ra.

For further information see [Table Jump Instructions](#).

Prerequisites

C or Zca must also be configured.

32-bit equivalent

No direct equivalent encoding exists.

Operation

```
//This is not SAIL, it's pseudo-code. The SAIL hasn't been written yet.

# target_address is temporary internal state, it doesn't represent a real register
# Mem is byte indexed
# "index8" is the field from the encoding, not "index" passed to the C.JT/C.JALT
  in the assembler
# which is formed below

if (OPCODE=="C.JALT") {
    index = index8 - 64;
} else {
    index = index8;
}

switch(XLEN) {
    32: table_address[XLEN-1:0] = JVT.base + (index<<2);
    64: table_address[XLEN-1:0] = JVT.base + (index<<3);
}

//check for debug mode entry, trigger with timing=0 and action=1, haltreq or step
if ((debug_trigger(table_address) && MCONTROL.timing==0 && MCONTROL.action==1) ||
    external_debug_haltreq() || DCSR.step==1) {
    DPC          = current_PC;
    DCSR.cause = DCSR.step==1 ? 4 : external_debug_haltreq() ? 3 : 2;
    enter_debug_mode();
}
//check for breakpoint trigger which takes an exception with timing=0
} else if ((debug_trigger(table_address) && MCONTROL.timing==0) ||
            !can_access_instruction_memory(table_address)) {
    MEPC  = current_PC;
    MTVAL = table_address;
    MCAUSE = debug_trigger(table_address) ? BREAKPOINT : INSTRUCTION_ACCESS_FAULT;
    take_exception();
} else {
    //access the jump table
    switch(XLEN) {
        32: LW target_address, InstMemory[table_address][XLEN-1:0];
        64: LD target_address, InstMemory[table_address][XLEN-1:0];
    }

    //don't use haltreq or step here, only check the addresses
    //check for table_address after reading if timing=1
    if (debug_trigger(table_address) && MCONTROL.timing==1 && MCONTROL.action==1) {
        DPC          = current_PC;
        DCSR.cause = 2;
        enter_debug_mode();
    } else if (debug_trigger(table_address) && MCONTROL.timing==1) {
```

```

    MEPC      = current_PC;
    MTVAL     = table_address;
    MCAUSE    = BREAKPOINT;
    take_exception();
} else if ((debug_trigger(target_address) && MCONTROL.timing==0 &&
MCONTROL.action==1) {
    DPC       = target_address;
    DCSR.cause = 2;
    enter_debug_mode();
} else if (((debug_trigger(target_address) && MCONTROL.timing==0) ||
            !can_access_instruction_memory(target_address)) {
    MEPC      = target_address;
    MTVAL     = target_address;
    MCAUSE    = debug_trigger(target_address) ? BREAKPOINT :
INSTRUCTION_ACCESS_FAULT;
    take_exception();
} else {
    //jump to the target address
    if (OPCODE=="C.JALT") {
        JALR ra, target_address[XLEN-1:0]&~0x1;
    } else {
        JR target_address[XLEN-1:0]&~0x1;
    }
}
}
}

```

Included in

Extension	Minimum version	Lifecycle state
Zcyt (Zcyt)	0.61.0	Stable

JVT CSR

Synopsis

Table jump base vector and control register


Address

TBD

Permissions

URW

Format (RV32, RV64)



Description

JVT.base is a virtual address, whenever virtual memory is enabled.

Using *JVT.base*[5:0] is implicitly zero, and is naturally aligned for all legal values of *XLEN*.

The memory pointed to by *JVT.base* is treated as instruction memory for the purpose of executing table jump instructions.

Table 2. *JVT.config* definition

JVT.config	Comment
000000	Jump table mode
others	reserved for future standard use

JVT.config is a WARL field, so can only be programmed to modes which are implemented. Therefore the discovery mechanism is to attempt to program different modes and read back the values to see which are available. Jump table mode *must* be implemented.

Architectural State

JVT adds architectural state to the context, therefore must be saved/restored on context switches.

Additional architectural state requires a state enable to be allocated. Accesses when the state is disabled will throw an illegal instruction exception. The state enable is not specified in this document.

Included in

Extension	Minimum version	Lifecycle state
Zcyt (Zcyt)	0.61.0	Stable

Table Jump Instructions

These instructions are collectively referred to as table jump:

- [c.jt](#): jump via table without link
- [c.jalt](#): jump via table and link to *ra*

Common details for these instructions are in this section.

Table Jump Overview

Table jump is a form of dictionary compression used to reduce the code size of JAL / AUIPC+JALR / JR / AUIPC+JR instructions.

Function calls and jumps to fixed labels typically take 32-bit or 64-bit instruction sequences.

Table jump allows the linker to:

- replace 32-bit *j* calls with *c.jt*
- replace 32-bit *jal ra* calls with *c.jalt*
- replace 64-bit *auipc/jalr* calls to fixed locations with *c.jt*
- replace 64-bit *auipc/jalr ra* calls to fixed locations with *c.jalt*
 - The AUIPC+JR/JALR sequence is used because the offset from the PC is out of the $\pm 1\text{MB}$ range.

JVT

The base of the table is in the JVT CSR (see [JVT CSR, table jump base vector and control register](#)), each table entry is XLEN bits.

The table entry number is from the *index8* field in the encoding, which controls the link register.

- *c.jt* : entries 0-63, link to *zero*
- *c.jalt* : entries 64-255, link to *ra*

Note that the LSB of every jump table entry is *ignored* which matches standard JALR behaviour.

If the same function is called with and without linking then it must have two entries in the table. This case does happen in practice but only affects a small number of entries so it does not waste much space in the table. It is typically caused by the same function being called with and without tail calling.

Recommended algorithm for allocating entries in the jump table

Calls to each function are categorised as shown in [Table jump code size saving for each function call replacement](#).

Table 3. Table jump code size saving for each function call replacement

original sequence	Table Jump saving
J	$A*2-(XLEN/8)$ bytes
AUIPC+JR	$B*6-(XLEN/8)$ bytes
JAL ra	$C*2-(XLEN/8)$ bytes
AUIPC+JALR ra	$D*6-(XLEN/8)$ bytes

Each function is called by using one of the two link registers. The total saving per function is calculated by counting the number of calls and adding up the total saving from each replacement of the existing sequence with a Table Jump instruction, as follows:

$$\begin{aligned} \text{saving_per_function_c_jt} &= A * 2 + B * 6 - 2*(XLEN-8) \\ \text{saving_per_function_c_jalt} &= C * 2 + D * 6 - 2*(XLEN-8) \end{aligned}$$

The functions are sorted so that the one with the highest saving is in table entry 0, the second highest in entry 1 etc. for that encoding.

NOTE

This algorithm assumes that each function is only called with one link register. If the same function is called with more than one link register, then it must have two entries in the table.

This allows the core to cache the most frequent targets by caching the lowest numbered entries of each section of the jump table. Only caching a few entries will greatly improve the performance.

Table Jump Fault handling

Table Jump involves two instruction fetches from a single instruction, and either fetch can cause a fault. There are no data accesses involved in the execution of table jumps.

The sequence required to execute the table jump instruction may be interrupted, or may not be able to start execution for several reasons.

- virtual memory page fault or PMP fault
 - these can be detected before execution, or during execution if the table jump instruction and the address in the table map to different virtual memory pages or PMP regions
- watchpoint trigger or debug halt caused by a trigger
- external debug halt
 - the halt can treat the whole sequence atomically, or interrupt mid sequence (implementation defined)
- interrupts
 - these may arrive at any time. An implementation can choose whether to interrupt the sequence or not.

For exceptions and interrupts MEPC contain the PC of the table jump instruction, MCAUSE is set as expected for the type of fault and MTVAL contains the address which caused the fault.

For debug halts DPC is set to the PC of the table jump instruction.

This section gives an overview of the behaviour, the exact operation is documented in the SAIL code for each instruction: [c.jalt SAIL code](#), [c.jt SAIL code](#).

Included in

Extension	Minimum version	Lifecycle state
Zcyp (Zcyp)	0.61.0	Stable