Zces 0.52

This document is in the Stable state. Assume anything could still change, but limited change should be expected. For more information see: https://riscv.org/spec-state

Zces is the set of sequenced or more complex instuctions for code-size reduction.

All 16-bit encodings are currently reserved for all architectures, and have no conflicts with any existing extensions.

All 32-bit encodings have yet to be allocated.

tblj and tbljal require tbljalvec CSR, table jump base vector and control register.

RV 32	RV 64	RV 128	Mnemonic	Instruction
✓	✓	✓	c.push {reg_list}, {areg_list}, -sp_adj	c.push: push registers to stack memory, 16-bit encoding
✓	✓	✓	push {reg_list}, {areg_list}, -sp_adj	push: push registers to stack memory, 32-bit encoding
✓	✓	✓	c.pop {reg_list}, sp_adjustment	c.pop: pop registers from the stack, 16-bit encoding
✓	✓	✓	pop {reg_list}, sp_adjustment	pop: pop registers from the stack, 32-bit encoding
✓	✓	✓	c.popret {reg_list}, {ret_val}, sp_adj	c.popret: pop registers and return, 16-bit encoding
✓	✓	✓	popret {reg_list}, {ret_val}, sp_adj	popret: pop registers from the stack and return, 32-bit encoding
✓	✓	✓	c.tblj #index	c.tblj: table jump without link, 16-bit encoding
✓	✓	✓	c.tbljal #index	c.tbljal: table jump and link to ra, 16-bit encoding
✓	✓	✓	c.mva01s07 sreg1, sreg2	c.mva01s07: move two s0-s7 registers into a0-a1, 16-bit encoding

c.push

Synopsis

Push registers to stack memory, 16-bit encoding

Mnemonic

c.push *{reg_list}, {areg_list}, -stack_adj*

Encoding (RV32, RV64, RV128)

15					10	9		7	6	5	4		2	1	0
1	0	0	0	1	1		spimm0_5[6:4]		1	0		rlist3		0	0
			FUNCT6									reg_list		OP=C0	

> **NOTE**   *spimm0_5* is only valid for *c.push* for values 0-5. Values 6 and 7 do *not* decode as *c.push*.

Syntax

```
c.push {<reg_list_16u> | <xreg_list_16u>}, {<areg_list>), -<stack_adj>
```

The variables used in the syntax are defined below.

```
<reg_list_16u> ::= <ra> ["," <s0>  | <s0-sN> ] (where N is 1,2,3,5,7,11)

if (<reg_list_16u>=="ra")           <xreg_list_16u>="x1"
if (<reg_list_16u>=="ra, s0")       <xreg_list_16u>="x1, x8"
if (<reg_list_16u>=="ra, s0-s1")    <xreg_list_16u>="x1, x8-x9"
if (<reg_list_16u>=="ra, s0-s2")    <xreg_list_16u>="x1, x8-x9, x18"
if (<reg_list_16u>=="ra, s0-sN")    <xreg_list_16u>="x1, x8-x9, x18-xM" (where
M=N+16 and N is 3,5,7,11)

if (<reg_list_16u>=="ra")           <areg_list>=""
if (<reg_list_16u>=="ra, s0")       <areg_list>="a0"
if (<reg_list_16u>=="ra, s0-sN")    <areg_list>="a0-aP" (where N is 1,2,3,5,7,11;
if (N<4) P=N; else P=3;)

if (<reg_list_16u>=="ra")           <stack_adj>=[16|32|48|64|96]
if (<reg_list_16u>=="ra, s0")       <stack_adj>=[16|32|48|64|96]
if (<reg_list_16u>=="ra, s0-s1")    <stack_adj>=[16|32|48|64|96]
if (<reg_list_16u>=="ra, s0-s2")    <stack_adj>=[16|32|48|64|96]
if (<reg_list_16u>=="ra, s0-s3")    <stack_adj>=[32|48|64|96|112]
if (<reg_list_16u>=="ra, s0-s5")    <stack_adj>=[32|48|64|96|112]
if (<reg_list_16u>=="ra, s0-s7")    <stack_adj>=[48|64|96|112|128]
if (<reg_list_16u>=="ra, s0-s11")   <stack_adj>=[64|96|112|128|144]
```

Description

This instruction pushes (stores) the registers in *reg_list* to stack memory, moves *areg_list* into similarly numbered *s* registers, and then adjusts the stack pointer by *-stack_adj*. For further information see

PUSH/POP Register Instructions.

Field decoding

The mapping from the *rlist3* and *spimm0_5* fields in the encoding are as shown below.

Table 1. rlist3 decoding

rlist3	reg_list_16u	stack_adj_base
0	ra	16
1	ra, s0	16
2	ra, s0-s1	16
3	ra, s0-s2	16
4	ra, s0-s3	32
5	ra, s0-s5	32
6	ra, s0-s7	48
7	ra, s0-s11	64

stack_adj_base covers enough 16-byte blocks of memory to cover the registers in *reg_list_16u*.

spimm_0_5 is used to allocate extra stack space in 16-byte blocks.

The total stack adjustment is calculated as shown.

```
stack_adj = stack_adj_base+spimm0_5[6:4]*16
```

NOTE | *spimm0_5* is only valid for values 0-5.

Prerequisites

The C-extension must also be configured.

32-bit equivalent

push: push registers to stack memory, 32-bit encoding

Operation

```
//This is not SAIL, it's pseudo-code. The SAIL hasn't been written yet.

//RV64/RV128 must have a 16-byte aligned sp
if (misa.MXL>=2 && sp[3:0]) {take_illegal_instruction_exception();}
//RV32I might be using the EABI (8-byte alignment) or UABI (16-byte alignment, so
in hardware we can only check for 8)
if (misa.MXL==1 && sp[2:0]) {take_illegal_instruction_exception();}

if (misa.MXL==1) {bytes=4;}
if (misa.MXL==2) {bytes=8;}
else             {bytes=16;}
addr=sp-bytes;
switch(bytes) {
  4:  asm("sw ra, 0(addr)");
  8:  asm("sd ra, 0(addr)");
  16: asm("sq ra, 0(addr)");
}
for(i=31;i>=0;i--)  {
  //if register i is in xreg_list
  if (xreg_list[i]) {
    addr-=bytes;
    switch(bytes) {
      4:  asm("sw s[i], 0(addr)");
      8:  asm("sd s[i], 0(addr)");
      16: asm("sq s[i], 0(addr)");
    }
  }
}
//The sequence must be uninterruptible from this point
if (areg_list[a0]) asm("mv s0, a0");
if (areg_list[a1]) asm("mv s1, a1");
if (areg_list[a2]) asm("mv s2, a2");
if (areg_list[a3]) asm("mv s3, a3");

sp+=stack_adjustment; //decrement
```

Assembly examples

```
c.push  {ra, s0-s5}, {a0-a3}, -64
```

Encoding: *rlist3=5, spimm0_5[6:4]=2*

Equivalent sequence:

```
sw  s5, -4(sp);
sw  s4, -8(sp);
sw  s3, -12(sp);
sw  s2, -16(sp);
sw  s1, -20(sp);
sw  s0, -24(sp);
sw  ra, -28(sp);
mv  s0, a0
mv  s1, a1
mv  s2, a2
mv  s3, a3
addi sp, sp, -64;
```

```
c.push {ra, s0-s1}, {a0-a1}, -32
```

Encoding: *rlist3=2, spimm0_5[6:4]=1*

Equivalent sequence:

```
sw  s1, -4(sp);
sw  s0, -8(sp);
sw  ra, -12(sp);
mv  s0, a0
mv  s1, a1
addi sp, sp, -32;
```

Included in

| Extension	Minimum version	Lifecycle state
Zces (Zces 0.52)	0.52	Stable

push

Synopsis

Push registers to stack memory, 32-bit encoding

Mnemonic

push {*reg_list*}, {*<areg_list>*}, *-stack_adj*

Encoding (RV32, RV64, RV128)

31		21	20 19		16	15 14		12 11		7	6		0
	TBD		areg	rlist_0_12		TBD	TBD		spimm[8:4]			TBD	
	FUNCT6			reg_list			funct3					opcode	

> **NOTE** | *rlist_0_12* is only valid for *push* for values 0-12. Values 13-15 map onto *different* encodings.

Syntax

```
push {<reg_list_32u> | <xreg_list_32u>}, {<areg_list>}, -<stack_adj>
```

The variables used in the syntax are defined below.

```
<reg_list_32u> ::= <ra> ["," <s0>  | <s0-sN> ] (where N is 1,2,...,11)

if (<reg_list_32u>=="ra")         <xreg_list_32u>="x1"
if (<reg_list_32u>=="ra, s0")     <xreg_list_32u>="x1, x8"
if (<reg_list_32u>=="ra, s0-s1")  <xreg_list_32u>="x1, x8-x9"
if (<reg_list_32u>=="ra, s0-s2")  <xreg_list_32u>="x1, x8-x9, x18"
if (<reg_list_32u>=="ra, s0-sN")  <xreg_list_32u>="x1, x8-x9, x18-xM" (where
M=N+16 and N is 3-11)

if (<reg_list_32u>=="ra")          <areg_list>=""
if (<reg_list_32u>=="ra, s0")      <areg_list>="" | "a0"
if (<reg_list_32u>=="ra, s0-sN")   <areg_list>="" | "a0-aP" (where N is 1-11; if
(N<4) P=N; else P=3;)

if (<reg_list_32u>=="ra")          <stack_adj>=[16|32|..|512]
if (<reg_list_32u>=="ra, s0")      <stack_adj>=[16|32|..|512]
if (<reg_list_32u>=="ra, s0-sN")   <stack_adj>=[16|32|..|512] (where N is 1,2)
if (<reg_list_32u>=="ra, s0-sN")   <stack_adj>=[32|48|..|528] (where N is 3,4,5,6)
if (<reg_list_32u>=="ra, s0-sN")   <stack_adj>=[48|64|..|544] (where N is
7,8,9,10)
if (<reg_list_32u>=="ra, s0-s11")  <stack_adj>=[64|96|..|560]
```

Description

This instruction pushes (stores) the registers in *reg_list* to stack memory, and then adjusts the stack pointer by *-stack_adj*. For further information see PUSH/POP Register Instructions.

Prerequisites

None

Field decoding

The mapping from the *rlist* and *spimm* fields in the encoding are as shown below.

Table 2. rlist decoding

rlist_0_12	reg_list_32u	stack_adj_base
0	ra	16
1	ra, s0	16
2	ra, s0-s1	16
3	ra, s0-s2	16
4	ra, s0-s3	32
5	ra, s0-s4	32
6	ra, s0-s5	32
7	ra, s0-s6	32
8	ra, s0-s7	48
9	ra, s0-s8	48
10	ra, s0-s9	48
11	ra, s0-s10	48
12	ra, s0-s11	64

stack_adj_base covers enough 16-byte blocks of memory to cover the registers in *reg_list_32u*. *spimm* is used to allocate extra stack space in 16-byte blocks. The total stack adjustment is calculated as shown.

```
stack_adj = stack_adj_base+spimm[8:4]*16
```

Table 3. areg_list decoding

rlist_0_12	areg_list	
	areg=0	areg=1
0	""	""
1	""	a0
2	""	a0-a1
3	""	a0-a2
4-12	""	a0-a3

Operation

```
//This is not SAIL, it's pseudo-code. The SAIL hasn't been written yet.

//RV64/RV128 must have a 16-byte aligned sp
if (misa.MXL>=2 && sp[3:0]) {take_illegal_instruction_exception();}
//RV32I might be using the EABI (8-byte alignment) or UABI (16-byte alignment, so
in hardware we can only check for 8)
if (misa.MXL==1 && sp[2:0]) {take_illegal_instruction_exception();}

if (misa.MXL==1) {bytes=4;}
if (misa.MXL==2) {bytes=8;}
else             {bytes=16;}
addr=sp-bytes;
switch(bytes) {
   4:  asm("sw ra, 0(addr)");
   8:  asm("sd ra, 0(addr)");
   16: asm("sq ra, 0(addr)");
}
for(i=31;i>=0;i--)  {
   //if register i is in xreg_list
   if (xreg_list[i]) {
     addr-=bytes;
     switch(bytes) {
        4:  asm("sw s[i], 0(addr)");
        8:  asm("sd s[i], 0(addr)");
        16: asm("sq s[i], 0(addr)");
     }
   }
}
//The sequence must be uninterruptible from this point
if (areg_list[a0]) asm("mv s0, a0");
if (areg_list[a1]) asm("mv s1, a1");
if (areg_list[a2]) asm("mv s2, a2");
if (areg_list[a3]) asm("mv s3, a3");

sp+=stack_adjustment; //decrement
```

Assembly examples

```
push  {ra, s0-s4}, {a0-a3}, -528
```

Encoding: *rlist*=5, *spimm[8:4]*=0x1f, *areg*=1

Equivalent sequence:

```
sw   s4, -4(sp);
sw   s3, -8(sp);
sw   s2, -12(sp);
sw   s1, -16(sp);
sw   s0, -20(sp);
sw   ra, -24(sp);
mv   s0, a0
mv   s1, a1
mv   s2, a2
mv   s3, a3
addi sp, sp, -528;
```

```
push {ra, s0-s3}, {}, -32
```

Encoding: *rlist3*=2, *spimm[8:4]*=1, *areg*=0

Equivalent sequence:

```
sw   s3, -4(sp);
sw   s2, -8(sp);
sw   s1, -12(sp);
sw   s0, -16(sp);
sw   ra, -20(sp);
addi sp, sp, -32;
```

Included in

Extension	Minimum version	Lifecycle state
Zces (Zces 0.52)	0.52	Stable

c.popret

Synopsis

Pop registers and return, 16-bit encoding

Mnemonic

c.popret *{reg_list}, {ret_val}, stack_adj*

Encoding (RV32, RV64, RV128)

15					10	9	7	6	5	4		2	1	0
1	0	0	0	1	1	spimm0_5[6:4]		1	ret0	rlist3			0	0

FUNCT6   reg_list   OP=C0

> **NOTE** *spimm0_5* is only valid for *c.popret* for values 0-5. Values 6 and 7 do *not* decode as *c.popret*.

Syntax

```
c.popret {<reg_list_16u> | <xreg_list_16u>}, <stack_adj>
```

The variables used in the syntax are defined below.

```
<reg_list_16u> ::= <ra> ["," <s0>  | <s0-sN> ] (where N is 1,2,3,5,7,11)

if (<reg_list_16u>=="ra")         <xreg_list_16u>="x1"
if (<reg_list_16u>=="ra, s0")     <xreg_list_16u>="x1, x8"
if (<reg_list_16u>=="ra, s0-s1")  <xreg_list_16u>="x1, x8-x9"
if (<reg_list_16u>=="ra, s0-s2")  <xreg_list_16u>="x1, x8-x9, x18"
if (<reg_list_16u>=="ra, s0-sN")  <xreg_list_16u>="x1, x8-x9, x18-xM" (where
M=N+16 and N is 3,5,7,11)

if (<reg_list_16u>=="ra")          <stack_adj>=[16|32|48|64|96]
if (<reg_list_16u>=="ra, s0")      <stack_adj>=[16|32|48|64|96]
if (<reg_list_16u>=="ra, s0-s1")   <stack_adj>=[16|32|48|64|96]
if (<reg_list_16u>=="ra, s0-s2")   <stack_adj>=[16|32|48|64|96]
if (<reg_list_16u>=="ra, s0-s3")   <stack_adj>=[32|48|64|96|112]
if (<reg_list_16u>=="ra, s0-s5")   <stack_adj>=[32|48|64|96|112]
if (<reg_list_16u>=="ra, s0-s7")   <stack_adj>=[48|64|96|112|128]
if (<reg_list_16u>=="ra, s0-s11")  <stack_adj>=[64|96|112|128|144]
```

Description

This instruction pop (loads) the registers in *reg_list* from stack memory, and then adjusts the stack pointer by *stack_adj*. For further information see PUSH/POP Register Instructions.

Field decoding

The mapping from the *rlist3* and *spimm0_5* fields in the encoding are as shown below.

Table 4. rlist3 decoding

rlist3	reg_list_16u	stack_adj_base
0	ra	16
1	ra, s0	16
2	ra, s0-s1	16
3	ra, s0-s2	16
4	ra, s0-s3	32
5	ra, s0-s5	32
6	ra, s0-s7	48
7	ra, s0-s11	64

stack_adj_base covers enough 16-byte blocks of memory to cover the registers in *reg_list_16u*. *spimm_0_5* is used to allocate extra stack space in 16-byte blocks. The total stack adjustment is calculated as shown.

```
stack_adj = stack_adj_base+spimm0_5[6:4]*16
```

NOTE   *spimm0_5* is only valid for values 0-5.

Prerequisites

The C-extension must also be configured.

32-bit equivalent

popret: pop registers from the stack and return, 32-bit encoding

Operation

```
//This is not SAIL, it's pseudo-code. The SAIL hasn't been written yet.

//RV64/RV128 must have a 16-byte aligned sp
if (misa.MXL>=2 && sp[3:0]) {take_illegal_instruction_exception();}
//RV32I might be using the EABI (8-byte alignment) or UABI (16-byte alignment, so
in hardware we can only check for 8)
if (misa.MXL==1 && sp[2:0]) {take_illegal_instruction_exception();}

if (misa.MXL==1) {bytes=4;}
if (misa.MXL==2) {bytes=8;}
else             {bytes=16;}
addr=sp+stack_adjustment-bytes;
switch(bytes) {
   4:  asm("lw ra, 0(addr)");
   8:  asm("ld ra, 0(addr)");
   16: asm("lq ra, 0(addr)");
}
for(i=31;i>=0;i--)  {
   //if register i is in xreg_list
   if (xreg_list[i]) {
     addr-=bytes;
     switch(bytes) {
        4:  asm("lw s[i], 0(addr)");
        8:  asm("ld s[i], 0(addr)");
        16: asm("lq s[i], 0(addr)");
     }
   }
}
if (ret_val) {
    switch(ret_val) {
       "0":  asm("li a0, 0");
    }
}
//The sequence must be uninterruptible from this point
sp+=stack_adjustment; //increment
asm("ret");
```

Assembly examples

```
c.popret   {ra, s0-s7}, {0}, 160
```

Encoding: *rlist3=6*, *spimm0_5[6:4]=7*, *ret0=1*

Equivalent sequence:

```
lw   s7, 156(sp);
lw   s6, 152(sp);
lw   s5, 148(sp);
lw   s4, 144(sp);
lw   s3, 140(sp);
lw   s2, 136(sp);
lw   s1, 132(sp);
lw   s0, 128(sp);
lw   ra, 124(sp);
li   a0, 0;
addi sp, sp, 160;
ret
```

```
c.popret   {ra, s0-s7}, {}, 160
```

Encoding: *rlist3=6*, *spimm0_5[6:4]=7*, *ret0=0*

Equivalent sequence:

```
lw   s7, 156(sp);
lw   s6, 152(sp);
lw   s5, 148(sp);
lw   s4, 144(sp);
lw   s3, 140(sp);
lw   s2, 136(sp);
lw   s1, 132(sp);
lw   s0, 128(sp);
lw   ra, 124(sp);
addi sp, sp, 160;
ret
```

Included in

| Extension	Minimum version	Lifecycle state
Zces (Zces 0.52)	0.52	Stable

popret

Synopsis

popret registers, 32-bit encoding

Mnemonic

popret *{reg_list}, {retval}, stack_adj*

Encoding (RV32, RV64, RV128)

31	22	21 20	19	16	15 14	12 11	7	6	0
TBD		ret_val	rlist_0_12		TBD	TBD	spimm[8:4]	TBD	
FUNCT6			reg_list			funct3		opcode	

> **NOTE**   *rlist_0_12* is only valid for *popret* for values 0-12. Values 13-15 map onto *different* encodings.

Syntax

```
popret {<reg_list_32u> | <xreg_list_32u>}, <stack_adj>
```

The variables used in the syntax are defined below.

```
<reg_list_32u> ::= <ra> ["," <s0>  | <s0-sN> ] (where N is 1,2,...,11)

if (<reg_list_32u>=="ra")        <xreg_list_32u>="x1"
if (<reg_list_32u>=="ra, s0")    <xreg_list_32u>="x1, x8"
if (<reg_list_32u>=="ra, s0-s1") <xreg_list_32u>="x1, x8-x9"
if (<reg_list_32u>=="ra, s0-s2") <xreg_list_32u>="x1, x8-x9, x18"
if (<reg_list_32u>=="ra, s0-sN") <xreg_list_32u>="x1, x8-x9, x18-xM" (where
M=N+16 and N is 3-11)

if (<reg_list_32u>=="ra")        <stack_adj>=[16|32|..|512]
if (<reg_list_32u>=="ra, s0")    <stack_adj>=[16|32|..|512]
if (<reg_list_32u>=="ra, s0-sN") <stack_adj>=[16|32|..|512] (where N is 1,2)
if (<reg_list_32u>=="ra, s0-sN") <stack_adj>=[32|48|..|528] (where N is 3,4,5,6)
if (<reg_list_32u>=="ra, s0-sN") <stack_adj>=[48|64|..|544] (where N is
7,8,9,10)
if (<reg_list_32u>=="ra, s0-s11") <stack_adj>=[64|96|..|560]
```

Description

This instruction pops (loads) the registers in *reg_list* from stack memory, and then adjusts the stack pointer by *stack_adj*. For further information see PUSH/POP Register Instructions.

Prerequisites

None

Field decoding

The mapping from the *ret_val*, *rlist* and *spimm* fields in the encoding are as shown below.

Table 5. rlist decoding

ret_val	reg_list_32u
0	no return value
1	a0=0
2	a0=1
3	a0=-1

Table 6. rlist decoding

rlist3	reg_list_32u	stack_adj_base
0	ra	16
1	ra, s0	16
2	ra, s0-s1	16
3	ra, s0-s2	16
4	ra, s0-s3	32
5	ra, s0-s4	32
6	ra, s0-s5	32
7	ra, s0-s6	32
8	ra, s0-s7	48
9	ra, s0-s8	48
10	ra, s0-s9	48
11	ra, s0-s10	48
12	ra, s0-s11	64

stack_adj_base covers enough 16-byte blocks of memory to cover the registers in *reg_list_32u*. *spimm* is used to allocate extra stack space in 16-byte blocks. The total stack adjustment is calculated as shown.

```
stack_adj = stack_adj_base+spimm[8:4]*16
```

Operation

```
//This is not SAIL, it's pseudo-code. The SAIL hasn't been written yet.


//RV64/RV128 must have a 16-byte aligned sp
if (misa.MXL>=2 && sp[3:0]) {take_illegal_instruction_exception();}
//RV32I might be using the EABI (8-byte alignment) or UABI (16-byte alignment, so
in hardware we can only check for 8)
if (misa.MXL==1 && sp[2:0]) {take_illegal_instruction_exception();}


if (misa.MXL==1) {bytes=4;}
if (misa.MXL==2) {bytes=8;}
else             {bytes=16;}
addr=sp+stack_adjustment-bytes;
switch(bytes) {
  4:  asm("lw ra, 0(addr)");
  8:  asm("ld ra, 0(addr)");
  16: asm("lq ra, 0(addr)");
}
for(i=31;i>=0;i--)  {
  //if register i is in xreg_list
  if (xreg_list[i]) {
    addr-=bytes;
    switch(bytes) {
      4:  asm("lw s[i], 0(addr)");
      8:  asm("ld s[i], 0(addr)");
      16: asm("lq s[i], 0(addr)");
    }
  }
}
if (ret_val) {
   switch(ret_val) {
      "0":  asm("li a0, 0");
      "1":  asm("li a0, 1");
      "2":  asm("li a0, -1");
   }
}
//The sequence must be uninterruptible from this point
sp+=stack_adjustment; //increment
asm("ret");
```

Assembly examples

```
popret    {ra, s0-s6}, {0}, 160
```

Encoding: *rlist=7*, *spimm[8:4]=7*, *ret0=1*

Equivalent sequence:

```
lw    s6, 156(sp);
lw    s5, 152(sp);
lw    s4, 148(sp);
lw    s3, 144(sp);
lw    s2, 140(sp);
lw    s1, 136(sp);
lw    s0, 132(sp);
lw    ra, 128(sp);
li    a0, 0;
addi sp, sp, 160;
ret
```

```
popret    {ra, s0-s7}, {-1}, 160
```

Encoding: *rlist=8*, *spimm[8:4]=7*, *ret0=2*

Equivalent sequence:

```
lw    s7, 156(sp);
lw    s6, 152(sp);
lw    s5, 148(sp);
lw    s4, 144(sp);
lw    s3, 140(sp);
lw    s2, 136(sp);
lw    s1, 132(sp);
lw    s0, 128(sp);
lw    ra, 124(sp);
li    a0, -1;
addi sp, sp, 160;
ret
```

Included in

Extension	Minimum version	Lifecycle state
Zces (Zces 0.52)	0.52	Stable

c.pop

Synopsis

Pop registers, 16-bit encoding

Mnemonic

c.pop *{reg_list}, stack_adj*

Encoding (RV32, RV64, RV128)

15							8	7	6	5	4	2	1	0
1	0	0	0	1	1	1	1	$pimm[4]	1	0	rlist3		0	0
			FUNCT6								reg_list		OP=C0	

Syntax

```
c.pop {<reg_list_16u> | <xreg_list_16u>}, <stack_adj>
```

The variables used in the syntax are defined below.

```
<reg_list_16u> ::= <ra> ["," <s0>  | <s0-sN> ] (where N is 1,2,3,5,7,11)

if (<reg_list_16u>=="ra")         <xreg_list_16u>="x1"
if (<reg_list_16u>=="ra, s0")     <xreg_list_16u>="x1, x8"
if (<reg_list_16u>=="ra, s0-s1")  <xreg_list_16u>="x1, x8-x9"
if (<reg_list_16u>=="ra, s0-s2")  <xreg_list_16u>="x1, x8-x9, x18"
if (<reg_list_16u>=="ra, s0-sN")  <xreg_list_16u>="x1, x8-x9, x18-xM" (where
M=N+16 and N is 3,5,7,11)

if (<reg_list_16u>=="ra")          <stack_adj>=[16|32]
if (<reg_list_16u>=="ra, s0")      <stack_adj>=[16|32]
if (<reg_list_16u>=="ra, s0-s1")   <stack_adj>=[16|32]
if (<reg_list_16u>=="ra, s0-s2")   <stack_adj>=[16|32]
if (<reg_list_16u>=="ra, s0-s3")   <stack_adj>=[32|48]
if (<reg_list_16u>=="ra, s0-s5")   <stack_adj>=[32|48]
if (<reg_list_16u>=="ra, s0-s7")   <stack_adj>=[48|64]
if (<reg_list_16u>=="ra, s0-s11")  <stack_adj>=[64|96]
```

Description

This instruction pop (loads) the registers in *reg_list* from stack memory, and then adjusts the stack pointer by *stack_adj*. For further information see PUSH/POP Register Instructions.

Field decoding

The mapping from the *rlist3* field in the encoding are as shown below.

Table 7. rlist3 decoding

rlist3	reg_list_16u	stack_adj_base
0	ra	16
1	ra, s0	16
2	ra, s0-s1	16
3	ra, s0-s2	16
4	ra, s0-s3	32
5	ra, s0-s5	32
6	ra, s0-s7	48
7	ra, s0-s11	64

stack_adj_base covers enough 16-byte blocks of memory to cover the registers in *reg_list_16u*. *spimm* is used to allocate extra stack space in 16-byte blocks. The total stack adjustment is calculated as shown.

```
stack_adj = stack_adj_base+spimm[4]*16
```

Prerequisites

The C-extension must also be configured.

32-bit equivalent

pop: pop registers from the stack, 32-bit encoding

Operation

```
//This is not SAIL, it's pseudo-code. The SAIL hasn't been written yet.

//RV64/RV128 must have a 16-byte aligned sp
if (misa.MXL>=2 && sp[3:0]) {take_illegal_instruction_exception();}
//RV32I might be using the EABI (8-byte alignment) or UABI (16-byte alignment, so
in hardware we can only check for 8)
if (misa.MXL==1 && sp[2:0]) {take_illegal_instruction_exception();}

if (misa.MXL==1) {bytes=4;}
if (misa.MXL==2) {bytes=8;}
else             {bytes=16;}
addr=sp+stack_adjustment-bytes;
switch(bytes) {
  4:  asm("lw ra, 0(addr)");
  8:  asm("ld ra, 0(addr)");
  16: asm("lq ra, 0(addr)");
}
for(i=31;i>=0;i--)  {
  //if register i is in xreg_list
  if (xreg_list[i]) {
    addr-=bytes;
    switch(bytes) {
      4:  asm("lw s[i], 0(addr)");
      8:  asm("ld s[i], 0(addr)");
      16: asm("lq s[i], 0(addr)");
    }
  }
}
//The sequence must be uninterruptible from this point
sp+=stack_adjustment; //increment
```

Assembly examples

```
c.pop   {ra, s0-s7}, 160
```

Encoding: *rlist3*=6, *spimm[4]*=0 Equivalent sequence:

```
lw   s7, 44(sp);
lw   s6, 40(sp);
lw   s5, 36(sp);
lw   s4, 32(sp);
lw   s3, 28(sp);
lw   s2, 24(sp);
lw   s1, 20(sp);
lw   s0, 16(sp);
lw   ra, 12(sp);
addi sp, sp, 48;
```

```
c.pop   {ra, s0-s7}, 160
```

Encoding: *rlist3*=6, *spimm[4]*=1

Equivalent sequence:

```
lw   s7, 60(sp);
lw   s6, 56(sp);
lw   s5, 52(sp);
lw   s4, 48(sp);
lw   s3, 44(sp);
lw   s2, 40(sp);
lw   s1, 36(sp);
lw   s0, 32(sp);
lw   ra, 28(sp);
addi sp, sp, 64;
```

Included in

Extension	Minimum version	Lifecycle state
Zces (Zces 0.52)	0.52	Stable

pop

Synopsis

Pop registers, 32-bit encoding

Mnemonic

pop {*reg_list*}, *stack_adj*

Encoding (RV32, RV64, RV128)

31	22	21 20	19	16	15 14	12 11	7 6	0
TBD		ret_val	rlist_0_12		TBD	TBD	spimm[8:4]	TBD
FUNCT6			reg_list			funct3		opcode
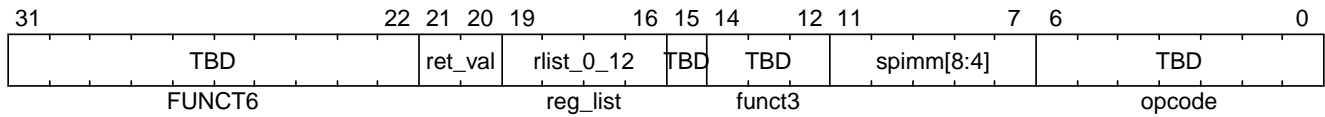
> **NOTE** | *rlist_0_12* is only valid for *pop* for values 0-12. Values 13-15 map onto *different* encodings.

Syntax

```
pop {<reg_list_32u> | <xreg_list_32u>}, <stack_adj>
```

The variables used in the syntax are defined below.

```
<reg_list_32u> ::= <ra> ["," <s0>  | <s0-sN> ] (where N is 1,2,...,11)

if (<reg_list_32u>=="ra")        <xreg_list_32u>="x1"
if (<reg_list_32u>=="ra, s0")    <xreg_list_32u>="x1, x8"
if (<reg_list_32u>=="ra, s0-s1") <xreg_list_32u>="x1, x8-x9"
if (<reg_list_32u>=="ra, s0-s2") <xreg_list_32u>="x1, x8-x9, x18"
if (<reg_list_32u>=="ra, s0-sN") <xreg_list_32u>="x1, x8-x9, x18-xM" (where
M=N+16 and N is 3-11)

if (<reg_list_32u>=="ra")        <stack_adj>=[16|32|..|512]
if (<reg_list_32u>=="ra, s0")    <stack_adj>=[16|32|..|512]
if (<reg_list_32u>=="ra, s0-sN") <stack_adj>=[16|32|..|512] (where N is 1,2)
if (<reg_list_32u>=="ra, s0-sN") <stack_adj>=[32|48|..|528] (where N is 3,4,5,6)
if (<reg_list_32u>=="ra, s0-sN") <stack_adj>=[48|64|..|544] (where N is
7,8,9,10)
if (<reg_list_32u>=="ra, s0-s11") <stack_adj>=[64|96|..|560]
```

Description

This instruction pops (loads) the registers in *reg_list* from stack memory, and then adjusts the stack pointer by *stack_adj*. For further information see PUSH/POP Register Instructions.

Prerequisites

None

Field decoding

The mapping from the *rlist* and *spimm* fields in the encoding are as shown below.

Table 8. rlist decoding

rlist3	reg_list_32u	stack_adj_base
0	ra	16
1	ra, s0	16
2	ra, s0-s1	16
3	ra, s0-s2	16
4	ra, s0-s3	32
5	ra, s0-s4	32
6	ra, s0-s5	32
7	ra, s0-s6	32
8	ra, s0-s7	48
9	ra, s0-s8	48
10	ra, s0-s9	48
11	ra, s0-s10	48
12	ra, s0-s11	64

stack_adj_base covers enough 16-byte blocks of memory to cover the registers in *reg_list_32u*. *spimm* is used to allocate extra stack space in 16-byte blocks. The total stack adjustment is calculated as shown.

```
stack_adj = stack_adj_base+spimm[8:4]*16
```

Prerequisites

The C-extension must also be configured.

Operation

```
//This is not SAIL, it's pseudo-code. The SAIL hasn't been written yet.

//RV64/RV128 must have a 16-byte aligned sp
if (misa.MXL>=2 && sp[3:0]) {take_illegal_instruction_exception();}
//RV32I might be using the EABI (8-byte alignment) or UABI (16-byte alignment, so
in hardware we can only check for 8)
if (misa.MXL==1 && sp[2:0]) {take_illegal_instruction_exception();}

if (misa.MXL==1) {bytes=4;}
if (misa.MXL==2) {bytes=8;}
else             {bytes=16;}
addr=sp+stack_adjustment-bytes;
switch(bytes) {
  4:  asm("lw ra, 0(addr)");
  8:  asm("ld ra, 0(addr)");
  16: asm("lq ra, 0(addr)");
}
for(i=31;i>=0;i--)  {
  //if register i is in xreg_list
  if (xreg_list[i]) {
    addr-=bytes;
    switch(bytes) {
      4:  asm("lw s[i], 0(addr)");
      8:  asm("ld s[i], 0(addr)");
      16: asm("lq s[i], 0(addr)");
    }
  }
}
//The sequence must be uninterruptible from this point
sp+=stack_adjustment; //increment
```

Assembly examples

```
pop   {ra, s0-s6}, 160
```

Encoding: *rlist=7, spimm[8:4]=7*

Equivalent sequence:

```
lw   s6, 156(sp);
lw   s5, 152(sp);
lw   s4, 148(sp);
lw   s3, 144(sp);
lw   s2, 140(sp);
lw   s1, 136(sp);
lw   s0, 132(sp);
lw   ra, 128(sp);
addi sp, sp, 160;
ret
```

```
pop   {ra, s0-s7}, 160
```

Encoding: *rlist=8, spimm[8:4]=7*

Equivalent sequence:

```
lw   s7, 156(sp);
lw   s6, 152(sp);
lw   s5, 148(sp);
lw   s4, 144(sp);
lw   s3, 140(sp);
lw   s2, 136(sp);
lw   s1, 132(sp);
lw   s0, 128(sp);
lw   ra, 124(sp);
addi sp, sp, 160;
ret
```

Included in

Extension	Minimum version	Lifecycle state
Zces (Zces 0.52)	0.52	Stable

c.tblj

Synopsis

table jump, no link, 16-bit encoding

Mnemonic

c.tblj *#index*

Encoding (RV32, RV64, RV128)

15				10	9		2	1	0
1	0	0	0	1	0	index8		0	0
		FUNCT6				index8<64		OP=C0	
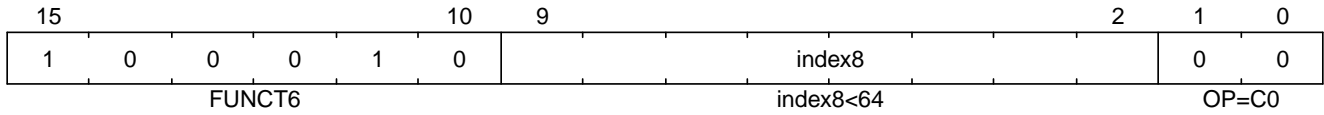
> **NOTE** For this encoding to decode as *TBLJ*, *index8<64*.

Syntax

```
c.tblj #index
```

Description

This instruction is used to dereference a table of PCs, and then jumps without linking to the dereferenced PC.

For further information see Table Jump Instructions.

Prerequisites

The C-extension must also be configured.

32-bit equivalent

No direct equivalent encoding exists.

Operation

```
//This is not SAIL, it's pseudo-code. The SAIL hasn't been written yet.

# target_address is temporary internal state, it doesn't represent a real register
# Mem is byte indexed
# index8 is the field from the encoding, not the index passed to the C.TBLJ* in
the assembler

switch(XLEN) {
  32:  table_address[XLEN-1:0] = TBLJALVEC.base + index8<<2;
  64:  table_address[XLEN-1:0] = TBLJALVEC.base + index8<<3;
  128: table_address[XLEN-1:0] = TBLJALVEC.base + index8<<4;
}


//check for debug mode entry, trigger with timing=0 and action=1, haltreq or step
if ((debug_trigger(table_address) && MCONTROL.timing==0 && MCONTROL.action==1) ||
    external_debug_haltreq() || DCSR.step==1) {
  DPC       = current_PC;
  DCSR.cause = DCSR.step==1 ? 4 : external_debug_haltreq() ? 3 : 2;
  enter_debug_mode();
//check for breakpoint trigger which takes an exception with timing=0
} else if ((debug_trigger(table_address) && MCONTROL.timing==0) ||
            !can_access_instruction_memory(table_address)) {
  MEPC   = current_PC;
  MTVAL  = table_address;
  MCAUSE = debug_trigger(table_address) ? BREAKPOINT : INSTRUCTION_ACCESS_FAULT;
  take_exception();
} else {
  //access the jump table
  switch(XLEN) {
    32:  LW target_address, InstMemory[table_address][XLEN-1:0];
    64:  LD target_address, InstMemory[table_address][XLEN-1:0];
    128: LQ target_address, InstMemory[table_address][XLEN-1:0];
  }

  //don't use haltreq or step here, only check the addresses
  //check for table_address after reading if timing=1
  if (debug_trigger(table_address) && MCONTROL.timing==1 && MCONTROL.action==1) {
    DPC       = current_PC;
    DCSR.cause = 2;
    enter_debug_mode();
  } else if (debug_trigger(table_address) && MCONTROL.timing==1) {
    MEPC      = current_PC;
    MTVAL     = table_address;
    MCAUSE    = BREAKPOINT;
    take_exception();
```

```
  } else if ((debug_trigger(target_address) && MCONTROL.timing==0 &&
MCONTROL.action==1) {
    DPC       = target_address;
    DCSR.cause = 2;
    enter_debug_mode();
  } else if (((debug_trigger(target_address) && MCONTROL.timing==0) ||
             !can_access_instruction_memory(target_address)) {
    MEPC      = target_address;
    MTVAL     = target_address;
    MCAUSE    = debug_trigger(target_address) ? BREAKPOINT :
INSTRUCTION_ACCESS_FAULT;
    take_exception();
  } else {
    //jump to the target address
    JALR zero, target_address[XLEN-1:0]&~0x1;
  }
}
```

Included in

Extension	Minimum version	Lifecycle state
Zces (Zces 0.52)	0.52	Stable

c.tbljal

Synopsis

table jump and link to ra, 16-bit encoding

Mnemonic

c.tbljal *#index*

Encoding (RV32, RV64, RV128)

15					10	9		2	1	0
1	0	0	0	1	0	index8			0	0
		FUNCT6				index8>=64			OP=C0	

> **NOTE** For this encoding to decode as *TBLJAL*, *index8>=64*.

Syntax

```
c.tbljal #index
```

Description

This instruction is used to dereference a table of PCs, and then jumps to the dereferenced PC and links to ra.

For further information see Table Jump Instructions.

Prerequisites

The C-extension must also be configured.

32-bit equivalent

No direct equivalent encoding exists.

Operation

```
//This is not SAIL, it's pseudo-code. The SAIL hasn't been written yet.

# target_address is temporary internal state, it doesn't represent a real register
# Mem is byte indexed
# index8 is the field from the encoding, not the index passed to the C.TBLJ* in
the assembler

switch(XLEN) {
  32:  table_address[XLEN-1:0] = TBLJALVEC.base + index8<<2;
  64:  table_address[XLEN-1:0] = TBLJALVEC.base + index8<<3;
  128: table_address[XLEN-1:0] = TBLJALVEC.base + index8<<4;
}


//check for debug mode entry, trigger with timing=0 and action=1, haltreq or step
if ((debug_trigger(table_address) && MCONTROL.timing==0 && MCONTROL.action==1) ||
    external_debug_haltreq() || DCSR.step==1) {
  DPC       = current_PC;
  DCSR.cause = DCSR.step==1 ? 4 : external_debug_haltreq() ? 3 : 2;
  enter_debug_mode();
//check for breakpoint trigger which takes an exception with timing=0
} else if ((debug_trigger(table_address) && MCONTROL.timing==0) ||
           !can_access_instruction_memory(table_address)) {
  MEPC   = current_PC;
  MTVAL  = table_address;
  MCAUSE = debug_trigger(table_address) ? BREAKPOINT : INSTRUCTION_ACCESS_FAULT;
  take_exception();
} else {
  //access the jump table
  switch(XLEN) {
    32:  LW target_address, InstMemory[table_address][XLEN-1:0];
    64:  LD target_address, InstMemory[table_address][XLEN-1:0];
    128: LQ target_address, InstMemory[table_address][XLEN-1:0];
  }

  //don't use haltreq or step here, only check the addresses
  //check for table_address after reading if timing=1
  if (debug_trigger(table_address) && MCONTROL.timing==1 && MCONTROL.action==1) {
    DPC       = current_PC;
    DCSR.cause = 2;
    enter_debug_mode();
  } else if (debug_trigger(table_address) && MCONTROL.timing==1) {
    MEPC      = current_PC;
    MTVAL     = table_address;
    MCAUSE    = BREAKPOINT;
    take_exception();
```

```
    } else if ((debug_trigger(target_address) && MCONTROL.timing==0 &&
  MCONTROL.action==1) {
      DPC        = target_address;
      DCSR.cause = 2;
      enter_debug_mode();
    } else if (((debug_trigger(target_address) && MCONTROL.timing==0) ||
                !can_access_instruction_memory(target_address)) {
      MEPC       = target_address;
      MTVAL      = target_address;
      MCAUSE     = debug_trigger(target_address) ? BREAKPOINT :
  INSTRUCTION_ACCESS_FAULT;
      take_exception();
    } else {
      //jump to the target address
      JALR ra, target_address[XLEN-1:0]&~0x1;
    }
  }
```

Included in

Extension	Minimum version	Lifecycle state
Zces (Zces 0.52)	0.52	Stable

TBLJALVEC CSR

Synopsis

Table jump base vector and control register

Address

TBD

Permissions

URW

Format (RV32, RV64, RV128)

31		6	5		0
	XLEN-1:6			config	
	base[XLEN-1:6]			CONFIGURATION	

Description

TBLJALVEC.base is a virtual address, whenever virtual memory is enabled.

Using *TBLJALVEC.base[5:0]* is implicitly zero, and is naturally aligned for all legal values of *XLEN*.

The memory pointed to by *TBLJALVEC.base* is treated as instruction memory for the purpose of executing table jump instructions.

Table 9. TBLJALVEC.config definition

TBLJALVEC.config	Comment
000000	Jump table mode
others	**reserved for future standard use**

TBLJALVEC.config is a WARL field, so can only be programmed to modes which are implemented. Therefore the discovery mechanism is to attempt to program different modes and read back the values to see which are available. Jump table mode *must* be implemented.

Architectural State

TBLJALVEC adds architectural state to the context, therefore must be saved/restore on context switch.

Additional architectural state requires a state enable to be allocated. Accesses when the state is disabled will throw an illegal instruction exception. The state enable is not specified in this document.

Included in

Extension	Minimum version	Lifecycle state
Zces (Zces 0.52)	0.52	Stable

c.mva01s07

Synopsis

Move two s0-s7 registers into a0-a1, 16-bit encoding

Mnemonic

c.mva01s07 *sreg1, sreg2*

Encoding (RV32, RV64, RV128)

15					10	9		7	6	5	4		2	1	0
0	1	0	0	1	1	1	sreg1		1	1	sreg2			0	1
		FUNCT6												OP=C1	

Syntax

```
c.mva01s07 sreg1, sreg2
```

Description

This instruction moves *sreg1* into *a0* and *sreg2* into *a1*. The execution is atomic, so it is not possible to observe state where only one of a0 or a1 have been updated.

Field decoding

Table 10. sreg decoding

sreg*	xreg
0	x8
1	x9
2	x18
3	x19
4	x20
5	x21
6	x22
7	x23

The encoding has two *sreg* number specifiers to save encoding space.

NOTE	This instruction does not directly expand to a single 32-bit encoding.

NOTE	The *s* register mapping is taken from the UABI, and may not match the currently unratified EABI.

Prerequisites

The C-extension must also be configured.

32-bit equivalent

No direct equivalent encoding exists.

Operation

```
//This is not SAIL, it's pseudo-code. The SAIL hasn't been written yet.


xreg1 = {sreg1[2:1]>0,sreg1[2:1]==0,sreg1[2:0]}
xreg2 = {sreg2[2:1]>0,sreg2[2:1]==0,sreg2[2:0]}


X[10] = X[sreg1]
X[11] = X[sreg2]
```

Included in

Extension	Minimum version	Lifecycle state
Zces (Zces 0.52)	0.52	Stable

PUSH/POP register instructions

These instructions are collectively referred to as PUSH/POP:

- c.push: push registers to stack memory, 16-bit encoding

- push: push registers to stack memory, 32-bit encoding

- c.popret: pop registers and return, 16-bit encoding

- popret: pop registers from the stack and return, 32-bit encoding

- c.pop: pop registers from the stack, 16-bit encoding

- pop: pop registers from the stack, 32-bit encoding

The term PUSH refers to both 16 and 32-bit encodings (C.PUSH, PUSH).

The term POP refers to both 16 and 32-bit encodings of POP (C.POP, POP).

The term POPRET refers to both 16 and 32-bit encodings of POPRET (C.POPRET, POPRET).

Common details for these instructions are in this section.

NOTE	This version does not fully include *areg_list* and will need to be updated if it's included in *push/c.push*.

PUSH/POP overview

PUSH, POP, POPRET along with the 16-bit forms are used to reduce the size of function prologues and epilogues.

1. The PUSH, C.PUSH instruction
    - pushes(stores) the registers specified in *reg_list* to the stack
    - if *areg_list* is included, moves the registers in the *areg_list* into *s* registers
        - *areg_list* is determined automatically from *rlist*, it cannot be arbitrarily specified. The definition is in c.push: push registers to stack memory, 16-bit encoding and push: push registers to stack memory, 32-bit encoding ;
    - adjusts the stack pointer by the *stack_adjustment* .

2. The POP, C.POP instruction .
    - pops(loads) the registers in *reg_list* from the stack.
    - adjusts the stack pointer by the *stack_adjustment*.

3. The POPRET, C.POPRET instruction
    - pops(loads) the registers in *reg_list* from the stack.
    - if *ret_val* is included, moves the specified constant value into *a0* as the return value.
    - adjusts the stack pointer by the *stack_adjustment*.
    - executes a *ret* instruction.

Example usage

This example gives an illustration of the use of PUSH and POPRET.

```
int function(void *buf, size_t len)
{
      return function2(buf, len);
}
```

compiles with GCC10 to:

```
20405458 <function>:
20405458: 1141                addi sp,sp,-16    ;#PUSH(1)
2040545a: c04a                sw   s2,0(sp)     ;#PUSH(2)
...
20405464: c422                sw   s0,8(sp)     ;#PUSH(3)
20405466: c226                sw   s1,4(sp)     ;#PUSH(4)
20405468: c606                sw   ra,12(sp)    ;#PUSH(5)
2040546a: 842a                mv   s0,a0        ;#PUSH(6)
2040546c: 84ae                mv   s1,a1        ;#PUSH(7)
<function body>
20405494: 4501                li   a0,0         ;#POPRET(1)
20405496: 40b2                lw   ra,12(sp)    ;#POPRET(2)
20405498: 4422                lw   s0,8(sp)     ;#POPRET(3)
2040549a: 4492                lw   s1,4(sp)     ;#POPRET(4)
2040549c: 4902                lw   s2,0(sp)     ;#POPRET(5)
2040549e: 0141                addi sp,sp,16     ;#POPRET(6)
204054a0: 8082                ret               ;#POPRET(7)
```

with the GCC option *-msave-restore* the output is the following:

```
204089ac <function>:
204089ac: f97f72ef            jal  t0,20400942 <__riscv_save_0> ;#PUSH(1)
...
204089b8: 842a                mv   s0,a0                        ;#PUSH(2)
204089ba: 84ae                mv   s1,a1                        ;#PUSH(3)
<function_body>
204089e2: 4501                li   a0,0                         ;#POPRET(1)
204089e4: f83f706f            j    20400966 <__riscv_restore_0> ;#POPRET(2)
```

with PUSH/POPRET this reduces to

```
20405458 <function>:
20405458: <16-bit>              push    {ra,s0-s2},{a0-a2},-16
<function body>
20405496: <16-bit>              popret {ra,s0-s2},{0}, 16
```

The prologue / epilogue reduce from 28-bytes in the original code, to 14-bytes with *-msave-restore*, and to 4-bytes with PUSH and POPRET. As well as reducing the code-size PUSH and POPRET eliminate the branches from calling the millicode *save/restore* routines so also perform better.

NOTE	The calls to *<riscv_save_0>/<riscv_restore_0>* become 64-bit when the target functions are out of the ±1MB range, increasing the prologue/epilogue size to 22-bytes.
NOTE	The C.PUSH has an additional register move included *mv s2, a2* which wasn't in the original prologue. This is included to simplify the encoding and definition of C.PUSH/PUSH and will cost some performance.
NOTE	POP is used for tail-calling which is not included in this example.

Compiler implementation

The technique used in the initial implementation in LLVM is to let the compiler generate the function prologue and epilogue, and then replace the instruction sequences with the relevant *push/pop* instructions.

reg_list handling

The actual saved/restored register list generated by the compiler may mismatch lists available in the 16-bit encodings. The number of occurances of this have been minimised during benchmarking, but cases will still occur.

When compiling with -Os/-Oz the compiler should add one or more registers and use the next shortest register list, even though this will cost performance and memory. Rounding up the register list is similar to the technique used for *-msave-restore* which saves/restores registers in groups of 4 only.

areg_list handling

c.push includes *areg_list*, and *push* optionally includes it.

Example: *c.push* fits perfectly

In this real world example generated by GCC10, *c.push* fits perfectly.

```
00e010b8 <function>:
  e010b8:        1141                    addi    sp,sp,-16 ; C.PUSH
  e010ba:        c422                    sw      s0,8(sp)  ; C.PUSH
  e010bc:        c226                    sw      s1,4(sp)  ; C.PUSH
  e010be:        c04a                    sw      s2,0(sp)  ; C.PUSH
  e010c0:        c606                    sw      ra,12(sp) ; C.PUSH
  e010c2:        842a                    mv      s0,a0     ; C.PUSH
  e010c4:        84ae                    mv      s1,a1     ; C.PUSH
  e010c6:        4908                    lw      a0,16(a0)
  e010c8:        4d8c                    lw      a1,24(a1)
  e010ca:        8932                    mv      s2,a2     ; C.PUSH
  e010cc:        726040ef                jal     ra,e057f2 <function2>
```

this is replaced by

```
00e010b8 <function1>:
  e010b8:        xxxx                    c.push {ra,s0-s2}, {a0-a2}, -16
  e010c6:        4908                    lw      a0,16(a0)
  e010c8:        4d8c                    lw      a1,24(a1)
  e010cc:        726040ef                jal     ra,e057f2 <function2>
```

Example: *reg_list* doesn't fit, *areg_list* doesn't fit

In this other real world example *areg_list* doesn't fit:

```
00e01126 <function3>:
  e01126:        1101                    addi    sp,sp,-32
  e01128:        ce06                    sw      ra,28(sp)
  e0112a:        cc22                    sw      s0,24(sp)
  e0112c:        ca26                    sw      s1,20(sp)
  e0112e:        c84a                    sw      s2,16(sp)
  e01130:        c64e                    sw      s3,12(sp)
  e01132:        c452                    sw      s4,8(sp)
  e01134:        c256                    sw      s5,4(sp)
  e01136:        c05a                    sw      s6,0(sp)
  e01138:        0e050363                beqz    a0,e0121e <function3+0xf8>
  e0113c:        8a2a                    mv      s4,a0
  e0113e:        852e                    mv      a0,a1
  e01140:        89ae                    mv      s3,a1
```

In this case, the required *reg_list* is not supported *and* the move instructions are not part of the same basic block, therefore compiling at -Os/-Oz would give:

```
00e01126 <function4>:
  e01126:        xxxx                     c.push {ra,s0-s7}, {a0-a3}, -32 ; include
saving s7 and moving {a0-a3} into {s0-s3}
  e01138:        0e050363                 beqz    a0,e0121e <function4+0xf8>
  e0113c:        8a2a                     mv      s4,a0
  e0113e:        852e                     mv      a0,a1
  e01140:        89ae                     mv      s3,a1
```

Compiling for performance would use *push* which perfectly fits the requirement but produces larger code:

```
00e01126 <function4>:
  e01126:        xxxxxxxx                 push {ra,s0-s6}, {}, -32
  e01138:        0e050363                 beqz    a0,e0121e <function4+0xf8>
  e0113c:        8a2a                     mv      s4,a0
  e0113e:        852e                     mv      a0,a1
  e01140:        89ae                     mv      s3,a1
```

Example: *areg_list* needs register allocation changes

The next case is where none of the register moves match the *areg_list* moves because the register allocator in the compiler did not allocate suitable registers:

```
00e01842 <function5>:

  e01e7e:        1101                     addi    sp,sp,-32
  e01e80:        cc22                     sw      s0,24(sp)
  e01e82:        c84a                     sw      s2,16(sp)
  e01e84:        c64e                     sw      s3,12(sp)
  e01e86:        c452                     sw      s4,8(sp)
  e01e88:        c256                     sw      s5,4(sp)
  e01e8a:        ce06                     sw      ra,28(sp)
  e01e8c:        ca26                     sw      s1,20(sp)
  e01e8e:        892a                     mv      s2,a0
  e01e90:        89ae                     mv      s3,a1
  e01e92:        8a32                     mv      s4,a2
  e01e94:        8ab6                     mv      s5,a3
  e01e96:        3f41                     jal     e01e26 <function6>
```

With *c.push* this becomes:

```
e01e7e <function5>:
  e01e7e:        1101                        c.push {ra,s0-s5}, {a0-a3}, -32; includes
moving {a0-a3} into {s0-s3}
  e01e8e:        892a                        mv      s2,a0; <-- switch dest to s0
  e01e90:        89ae                        mv      s3,a1; <-- switch dest to s1
  e01e92:        8a32                        mv      s4,a2; <-- switch dest to s2
  e01e94:        8ab6                        mv      s5,a3; <-- switch dest to s3
  e01e96:        3f41                        jal     e01e26 <function6>
```

In this case all four moves can be deleted if the register allocation can be altered.

Example: *areg_list* partially fits

In this final case, one register move can be deleted and one must be retained unless the register allocation can be changed.

```
00e02368 <function7>:
  e02368:        1141                        addi    sp,sp,-16
  e0236a:        c226                        sw      s1,4(sp)
  e0236c:        03450493                    addi    s1,a0,52
  e02370:        c422                        sw      s0,8(sp)
  e02372:        842a                        mv      s0,a0; <-- matches areg_list -
delete
  e02374:        8526                        mv      a0,s1; <-- doesn't fit areg_list
  e02376:        c04a                        sw      s2,0(sp)
  e02378:        c606                        sw      ra,12(sp)
  e0237a:        892e                        mv      s2,a1; <-- switch dest to s1
  e0237c:        df3fd0ef                    jal     ra,e0016e <function8>
```

```
00e02368 <function7>:
  e02368:        xxxx                        c.push {ra,s0-s2}, {a0-a2}, -16
  e0236c:        03450493                    addi    s1,a0,52
  e02374:        8526                        mv      a0,s1; <-- doesn't fit areg_list
  e0237a:        892e                        mv      s2,a1; <-- switch dest to s1
  e0237c:        df3fd0ef                    jal     ra,e0016e <function8>
```

In this case one move is deleted, but one remains because unless the target register can be reallocated.

For the smallest code-size the compiler should reallocate the target registers so that the moves in *areg_list* are not wasted.

Compiling PUSH/POP for size or performance

As mentioned above, there are cases where there are choices about whether to select the 16-bit or 32-bit encoding. The 32-bit encodings offer a smaller stack adjustment range than using a 16-bit encoding and an

additional C.ADDI16SP instruction. Therefore using the 32-bit encoding will not reduce the code size if the stack adjustment is out of range of the 16-bit encoding.

The main performance/code-size trade-offs are

- whether *reg_list* is available in the 16-bit encodings matches the required list, and so whether extra registers are included by the 16-bit encoding
- whether *areg_list* includes redundant moves

The recommendation is that the 32-bit encoding should be selected only if compiling for performance and either

- *reg_list* is not available in the 16-bit encoding
- *areg_list* includes redundant moves

In addition, for POPRET, the 32-bit encoding allows more return values than the 16-bit encoding. Therefore the recommendation is that the 32-bit encoding should be selected if the 32-bit encoding allows the required return value.

PUSH/POP Fault handling

The sequence required to execute the PUSH/POP instruction may be interrupted, or may not be able to start execution for several reasons.

- virtual memory page fault or PMP fault
  - these can be detected before execution, or during execution if the memory addresses cross a page/PMP boundary
  - MTVAL is set to any address which causes the fault
- watchpoint trigger
  - these can be detected before execution, or during execution depending on the trigger type (load data triggers require the sequence to have started executing, for example)
  - MTVAL is set to any address which causes the fault
- external debug halt
  - the halt can treat the whole sequence atomically, or interrupt mid sequence (implementation defined)
- debug halt caused by a trigger
  - same comment as watchpoint trigger above
- load access fault
  - these are detected while the sequence is executing
  - MTVAL is set to the fault address.
- store access fault (precise or imprecise)
  - these may be detected while the sequence is executing, or afterwards if imprecise
  - MTVAL is set to the fault address.
- interrupts
  - these may arrive at any time. An implementation can choose whether to interrupt the sequence or not.

NOTE	MTVAL may be hardwired to zero in an implementation. The section above assumes it is implemented.

In all case MEPC contain the PC of the PUSH/POP instruction, and MCAUSE is set as expected for the type of fault.

For debug halts DPC is set to the PC of the PUSH/POP instruction.

Because some faults can only be detected during the sequence the core implementation is able to recover from the fault and re-execute the sequence. This may involve executing some or all of the loads and stores from the sequence multiple times before the sequence completes (as multiple faults or multiple interrupts are possible).

Therefore correct execution requires that *sp* refers to idempotent memory (also see Non-idempotent memory handling).

Software view of execution

Software view of the PUSH sequence

From a software perspective the PUSH sequence appears as:

- A sequence of stores writing a contiguous block of memory. Any of the bytes may be written multiple times.
- A stack pointer adjustment

Because the memory is idempotent and the stores are non-overlapping, they may be reordered, grouped into larger accesses, split into smaller access or any combination of these.

If an implementation allows interrupts during the sequence, and the interrupt handler uses *sp* to allocate stack memory, then any stores which were executed before the interrupt may be overwritten by the handler. This is safe because the memory is idempotent and the stores will be re-executed execution resumes.

The stack pointer adjustment must only be committed once it is certain that all of the stores will complete within triggerring any precise faults (stores may return imprecise bus errors which are received after the instruction has completed execution).

For example:

```
c.push  {ra, s0-s5}, {a0-a3}, -64
```

Appears to software as:

```
# any bytes from sp-1 to sp-28 may be written multiple times before the
instruction completes
sw  s5, -4(sp);
sw  s4, -8(sp);
sw  s3,-12(sp);
sw  s2,-16(sp);
sw  s1,-20(sp);
sw  s0,-24(sp);
sw  ra,-28(sp);

# these must only execute once, and will only execute after all stores complete
sucessfully
mv   s0, a0
mv   s1, a1
mv   s2, a2
mv   s3, a3
addi sp, sp, -64;
```

Software view of the POP/POPRET sequence

From a software perspective the POP/POPRET sequence appears as:

- A sequence of loads, any of which may be executed multiple times
- A stack pointer adjustment
- An optional LI into a0
- An optional RET

If an implementation allows interrupts during the sequence, then any loads which were executed before the interrupt may update architectural state. The loads will be re-executed once the handler completes, so the values will be overwritten. Therefore it is permitted for an implementation to update some of the destination registers before taking the interrupt or other fault.

The load immediate and stack pointer adjustment must only be committed once it is certain that all of the loads will complete successfully.

For POPRET once the stack pointer adjustment has been committed the RET must execute.

For example:

```
popret   {ra, s0-s3}, {1}, 32 ;
```

Appears to software as:

```
# any or all of these load instructions may execute multiple times
lw    s3, 28(sp);
lw    s2, 24(sp);
lw    s1, 20(sp);
lw    s0, 16(sp);
lw    ra, 12(sp);


# must only execute once, will only execute after all loads complete successfully
# all instructions must execute atomically
li a0, 1
addi sp, sp, 32;
ret;
```

Non-idempotent memory handling

An implementation may have a requirement to issue a PUSH/POP instruction to non-idempotent memory.

Error detection

If the core implementation does not have a requirement to support PUSH/POP to non-idempotent memories, and the core can use a PMA to detect that the memory is non-idempotent, then take a load(POP/POPRET) or store (PUSH) access fault exception.

Non-idempotent support

It is possible to support non-idempotent memory. One reason is to re-use PUSH/POP as a restricted form of a load/store multiple instruction to a peripheral, as there is no generic load/store multiple instruction in the RISC-V ISA.

If accessing non-idempotent memory then it is *recommended* to:

1. Not allow interrupts during execution

2. Not allow external debug halt during execution

3. Detect any virtual memory page faults or PMP faults for the whole instruction before starting execution (instead of during the sequence)

4. Not split / merge / reorder the generated memory accesses

It is possible that one of the following will still occur during execution:

1. Watchpoint trigger

2. Load/store access fault

In these cases the core will jump to the debug or exception handler. If execution is required to continue afterwards (so the event is not fatal to the code execution), then the handler is required to do so in software.

By following these rules memory accesses will only ever be issued once, and in the order listed in the SAIL.

It is possible for implementations to follow these restricted rules and to safely access both types of memory. It is also possible for an implementation to use PMAs to detect the memory type and apply different rules, such as

only allowing interrupts if accessing cacheable memory, for example.

Included in

Extension	Minimum version	Lifecycle state
Zces ([zces])	0.52	Stable

Table Jump Instructions

These instructions are collectively referred to as table jump:

- c.tblj: table jump without link, 16-bit encoding
- c.tbljal: table jump and link to ra, 16-bit encoding

Common details for these instructions are in this section.

Table Jump Overview

Table jump is a form of dictionary compression used to reduce the code size of JAL / AUIPC+JALR / JR / AUIPC+JR instructions.

Function calls and jumps to fixed labels typically take 32-bit or 64-bit instruction sequences.

Table jump allows the linker to:

- replace 32-bit J calls with C.TBLJ
- replace 32-bit JAL ra calls with C.TBLJAL
- replace 64-bit AUIPC/JR calls to fixed locations with C.TBLJ
- replace 64-bit AUIPC/JALR ra calls to fixed locations with C.TBLJAL
    - The AUIPC+JR/JALR sequence is used because the offset from the PC is out of the $\pm$1MB range.

TBLJALVEC

The base of the table is in the TBLJALVEC CSR (see tbljalvec CSR, table jump base vector and control register), each table entry is XLEN bits.

The table entry number is from the *index8* field in the encoding, which controls the link register.

- C.TBLJ : entries 0-63, link to *zero*
- C.TBLJAL : entries 64-255, link to *ra*

Note that the LSB of every jump table entry is *ignored* which matches standard JALR behaviour.

If the same function is called with and without linking then it must have two entries in the table. This case does happen in practice but only affects a small number of entries so it does not waste much space in the table. It is typically caused by the same function being called with and without tail calling.

Recommended algorithm for allocating entries in the jump table

Calls to each function are categorised as shown in Table jump code size saving for each function call replacement.

Table 11. Table jump code size saving for each function call replacement

original sequence	Table Jump saving
J	A*2-(XLEN/8) bytes

original sequence	Table Jump saving
AUIPC+JR	B*6-(XLEN/8) bytes
JAL ra	C*2-(XLEN/8) bytes
AUIPC+JALR ra	D*6-(XLEN/8) bytes

Each function is called by using one of the two link registers. The total saving per function is calculated by counting the number of calls and adding up the total saving from each replacement of the existing sequence with a Table Jump instruction, as follows:

```
saving_per_function_c_tblj    = A * 2 + B * 6 - 2*(XLEN-8)
saving_per_function_c_tbljal  = C * 2 + D * 6 - 2*(XLEN-8)
```

The functions are sorted so that the one with the highest saving is in table entry 0, the second highest in entry 1 etc. for that encoding.

NOTE	This algorithm assumes that each function is only called with one link register. If the same function is called with more than one link register, then it must have two entries in the table.

This allows the core to cache the most frequent targets by caching the lowest numbered entries of each section of the jump table. Only caching a few entries will greatly improve the performance.

Table Jump Fault handling

Table Jump involves two instruction fetches from a single instruction, and either fetch can cause a fault.

The sequence required to execute the table jump instruction may be interrupted, or may not be able to start execution for several reasons.

- virtual memory page fault or PMP fault
  - these can be detected before execution, or during execution if the memory addresses cross a page/PMP boundary
  - MTVAL is set to any address which causes the fault
- watchpoint trigger
  - these can be detected before execution, or during execution depending on the trigger type (load data triggers require the sequence to have started executing, for example)
  - MTVAL is set to any address which causes the fault
- external debug halt
  - the halt can treat the whole sequence atomically, or interrupt mid sequence (implementation defined)
- debug halt caused by a trigger
  - same comment as watchpoint trigger above
- load access fault
  - these are detected while the sequence is executing
  - MTVAL is set to the fault address.

- store access fault (precise or imprecise)

  ○ these may be detected while the sequence is executing, or afterwards if imprecise

  ○ MTVAL is set to the fault address.

- interrupts

  ○ these may arrive at any time. An implementation can choose whether to interrupt the sequence or not.

In all case MEPC contain the PC of the table jump instruction, and MCAUSE is set as expected for the type of fault.

For debug halts DPC is set to the PC of the table jump instruction.

This section gives an overview of the behaviour, the exact operation is documented in the SAIL code for each instruction

- c.tbljal SAIL code

- c.tblj SAIL code

Included in

Extension
Minimum version
Lifecycle state
Zces ([zces])
0.52
Stable